



NEW YORK UNIVERSITY

Energy-Based Models, Variational Methods

Yann LeCun

NYU - Courant Institute & Center for Data Science

Facebook AI Research

<http://yann.lecun.com>

Deep Learning, NYU, Spring 2023

EBM & Probabilistic Models



Refresher on turning energies to probabilities

- ▶ Gibbs distribution (a.k.a. softmax, should be called softargmax)

- ▶ Discrete / Continuous

$$P_w(y) = \frac{e^{-\beta F_w(y)}}{\sum_{y'} e^{-\beta F_w(y')}} \quad P_w(y) = \frac{e^{-\beta F_w(y)}}{\int_{y'} e^{-\beta F_w(y')}}$$

- ▶ Joint distribution

$$P_w(y, z) = \frac{e^{-\beta E_w(y, z)}}{\int_{y'} \int_{z'} e^{-\beta E_w(y', z')}}$$

Partition function Inverse temperature

- ▶ Conditional distribution

$$P_w(y, z|x) = \frac{e^{-\beta E_w(x, y, z)}}{\int_{y'} \int_{z'} e^{-\beta E_w(x, y', z')}}$$

- ▶ Marginal distribution

$$P_w(y|x) = \int_{z'} P_w(y, z'|x) = \frac{\int_{z'} e^{-\beta E_w(x, y, z')}}{\int_{y'} \int_{z'} e^{-\beta E_w(x, y', z')}}$$

Refresher on turning energies to probabilities

► **Joint distribution**

$$P_w(y, z) = \frac{e^{-\beta E_w(y, z)}}{\int_{y'} \int_{z'} e^{-\beta E_w(y', z')}$$

► **Conditional distribution**

$$P_w(y|z) = \frac{e^{-\beta E_w(y, z)}}{\int_{y'} e^{-\beta E_w(y', z)}}$$

► **Marginal distribution**

$$P_w(z) = \frac{\int_{y'} e^{-\beta E_w(y', z)}}{\int_{z'} \int_{y'} e^{-\beta E_w(y', z')}$$

► **Bayes rules!**

$$P_w(y, z) = P_w(y|z)P_w(z) = P_w(z|y)P_w(y)$$

Negative log-likelihood loss

$$L(x, y, w) = -\frac{1}{\beta} \log P_w(y|x) = F_w(x, y) + \frac{1}{\beta} \log \left[\int_{y'} e^{-\beta F_w(x, y')} \right]$$

► **Gradient of log partition function**

Minus log partition function.
Like a free energy over y .

$$\frac{\partial \left[-\frac{1}{\beta} \log \left[\int_{y'} e^{-\beta F_w(x, y')} \right] \right]}{\partial w} = \int_{y'} P_w(y'|x) \frac{\partial F_w(x, y')}{\partial w}$$

► **Monte Carlo methods: sample y from $P(y|x)$**

- The integral is an expectation of the gradient over the distribution of y
- Sample y from the distribution and average the corresponding gradients.

Max Likelihood is (generally) a (bad) Contrastive Method

- ▶ Push down on data points,
- ▶ Push up on all points
- ▶ Max likelihood / probabilistic models

$$P_w(y|x) = \frac{e^{-\beta F_w(x,y)}}{\int_{y'} e^{-\beta F_w(x,y')}}$$

- ▶ Loss: $\mathcal{L}(x, y, w) = F_w(x, y) + \frac{1}{\beta} \log \int_{y'} e^{-\beta F_w(x,y')}$

- ▶ Gradient: $\frac{\partial \mathcal{L}(x, y, w)}{\partial w} = \frac{\partial F_w(x, y)}{\partial w} - \int_{y'} P_w(y'|x) \frac{\partial F_w(x, y')}{\partial w}$

- ▶ 2nd term is intractable: MC/MCMC/HMC/CD: \hat{y} sampled from $P_w(y|x)$

$$\frac{\partial \mathcal{L}(x, y, w)}{\partial w} = \frac{\partial F_w(x, y)}{\partial w} - \frac{\partial F_w(x, \hat{y})}{\partial w}$$

push down of the energy of data points, push up everywhere else

Gradient of the negative log-likelihood loss for one sample Y :


$$\frac{\partial \mathcal{L}(x, y, w)}{\partial w} = \frac{\partial F_w(x, y)}{\partial w} - \int_{y'} P_w(y'|x) \frac{\partial F_w(x, y')}{\partial w}$$

Gradient descent:

$$w \leftarrow w - \eta \frac{\partial \mathcal{L}(x, y, w)}{\partial w}$$

Pushes down on the energy of the samples


Pulls up on the energy of low-energy Y 's



$$w \leftarrow w - \eta \frac{\partial F_w(x, y)}{\partial w} + \eta \int_{y'} P_w(y'|x) \frac{\partial F_w(x, y')}{\partial w}$$


Problem with Max Likelihood / Probabilistic Methods

- ▶ It wants to make the difference between the energy on the data manifold and the energy just outside of it infinitely large!
- ▶ It wants to make the data manifold an infinitely deep and infinitely narrow canyon.
- ▶ The loss must be **regularized** to keep the energy smooth
 - ▶ e.g. with Bayesian prior or by limiting weight sizes à la Wasserstein GAN.
 - ▶ So that gradient-based inference works
 - ▶ Equivalent to a Bayesian prior
 - ▶ But then, why use a probabilistic model?



Regularization through (variational) marginalization.

Push down on the energy of training samples.
Minimize the capacity of the latent variables.
Maximize the capacity of the representation.



Making z a noisy variable to reduce its information content

- ▶ The information content of the latent variable z must be minimized


- ▶ One (probabilistic) way to do this:

$$E(y, z) = C(y, Dec(z))$$

- ▶ make z “fuzzy” (e.g. stochastic)
- ▶ Z is a sample from a distribution $q(z|y)$

- ▶ Minimize the expected value of the energy under $q(z|y)$

$$\langle E(y) \rangle = \int_z q(z|y) E(y, z)$$




- ▶ Minimize the information content of $q(z|y)$ about y

What information does $q(z|y)$ give us about y ?

- ▶ Suppose that all the z come from a distribution $p(z)$
 - ▶ e.g. $p(z)$ uniform over a hypercube of dimension d : $[-r/2, +r/2]^d$
- ▶ Suppose that $q(z|y)$ is uniform
 - ▶ over a small hypercube of size s centered on $m(y)$
 - ▶ e.g. $q(z|y)$ uniform over $[m_i(y)-s/2, m_i(y)+s/2]$ in each dimension i .
- ▶ There are $(r/s)^d$ small cubes in the big cube
 - ▶ Hence each small cube gives

$$H(z|y) = \log_2(r/s)^d = d \log_2(r/s)$$
 bits of information about y .
- ▶ To minimize the information content of z ...
 - ▶ I can make the small cube large
 - ▶ I can make the large cube small



What information does $q(z|y)$ give us about y ?


- ▶ Suppose that all the z come from a distribution $p(z)$
- ▶ Suppose that each z distributes according to $q(z|y)$
- ▶ The amount of information that $q(z|y)$ gives about $p(z)$ is

$$KL(q(z|y), p(z)) = \int_z q(z|y) \log_2(q(z|y)/p(z))$$

- ▶ **Example: uniform case:** $p(z) = (1/r)^d$, $q(z|y) = (1/s)^d$

$$KL(q(z|y), p(z)) = \int_z q(z|y) \log_2((1/s)^d / (1/r)^d)$$

$$KL(q(z|y), p(z)) = \log_2(r/s)^d \int_z q(z|y) = d \log_2(r/s)$$



General case: minimize expected energy & information of z on y

► Minimize the expected energy

$$\langle E(y) \rangle_q = \int_z q(z|y) E(y, z)$$


$$E(y, z) = C(y, Dec(z))$$

► Minimize the relative entropy

- Between $q(z|y)$ and a prior distribution $p(z)$.


$$KL(q(z|y), p(z)) = \int_z q(z|y) \log_2(q(z|y)/p(z))$$

- This is the number of bits one sample from $q(z|y)$ will give us about $p(z)$



Marginalization as Regularization through Maximum Entropy

- ▶ Find a distribution $q(z|y)$ that minimizes the expected energy while having maximum entropy
- ▶ high entropy distribution == small information content from a sample



- ▶ Pick a family of distributions $q(z|y)$ (e.g. Gaussians) and find the one that minimizes the **variational free energy**:

$$\tilde{F}_q(y) = \int_z q(z|y) E(y, z) + \frac{1}{\beta} \int_z q(z|y) \log_2(q(z|y)/p(z))$$

- ▶ The trade-off between energy and entropy is controlled by the beta parameter.

Gaussian case

- Both $p(z)$ and $q(z|y)$ are Gaussians

$$p(z) = N(0, r) \qquad q(z|y) = N(m(y), s(y))$$


$$KL(q, p) = \log(r/s(y)) + \frac{m(y)^2 + s(y)^2}{2r^2} - \frac{1}{2}$$

(this is in nats, not bits. Divide by $\log(2)$ to get it in bits).

- Assume $r=1$:

$$KL(q, p) = \frac{1}{2}(m(y)^2 + s(y)^2 - \log_2(s(y)^2) - 1)$$

- This has a minimum at $s=1$



Marginalization as Regularization through Maximum Entropy

$$\tilde{F}_q(y) = \int_z q(z|y) E(y, z) + \frac{1}{\beta} \int_z q(z|y) \log(q(z|y)/p(z))$$

- ▶ If the family $q(z|y)$ is flexible enough, the $q^*(z|y)$ that minimizes the variational free energy is the Gibbs distribution:

$$q^*(z|y) = \frac{e^{-\beta E(y, z)}}{\int_{z'} e^{-\beta E(y, z')}}$$

- ▶ With this $q^*(z|y)$, the variational free energy becomes the free energy:

$$\tilde{F}_{q^*}(y) = F(y) = -\frac{1}{\beta} \log \int_z e^{-\beta E(y, z)}$$

- ▶ The Gibbs distribution on z is the one best trade-off between minimizing the expected energy and maximizing its entropy

Variational Inference

► Variational approximation of marginalization over z

$$F(y) = -\frac{1}{\beta} \log \int_z q(z|y) \frac{e^{-\beta E(y,z)}}{q(z|y)}$$


Jensen's inequality: $-\log(\text{average}()) < \text{average}(-\log())$

$$F(y) \leq \tilde{F}(y) = \int_z q(z|y) \left[-\frac{1}{\beta} \log \frac{e^{-\beta E(y,z)}}{q(z|y)} \right]$$

► Variational free energy:

$$\tilde{F}(y) = \int_z q(z|y) E(y, z) + \frac{1}{\beta} \int_z q(z|y) \log(q(z|y))$$

$F = \langle E \rangle - TS$ ← Helmholtz free energy in thermodynamics:
for a given average energy $\langle E \rangle$, a system minimizes its free energy by maximizing its entropy S . The trade-off depends on the temperature T



Variational Auto-Encoder

- If $Q(y,z)$ is quadratic, $q(z|y)$ is Gaussian.

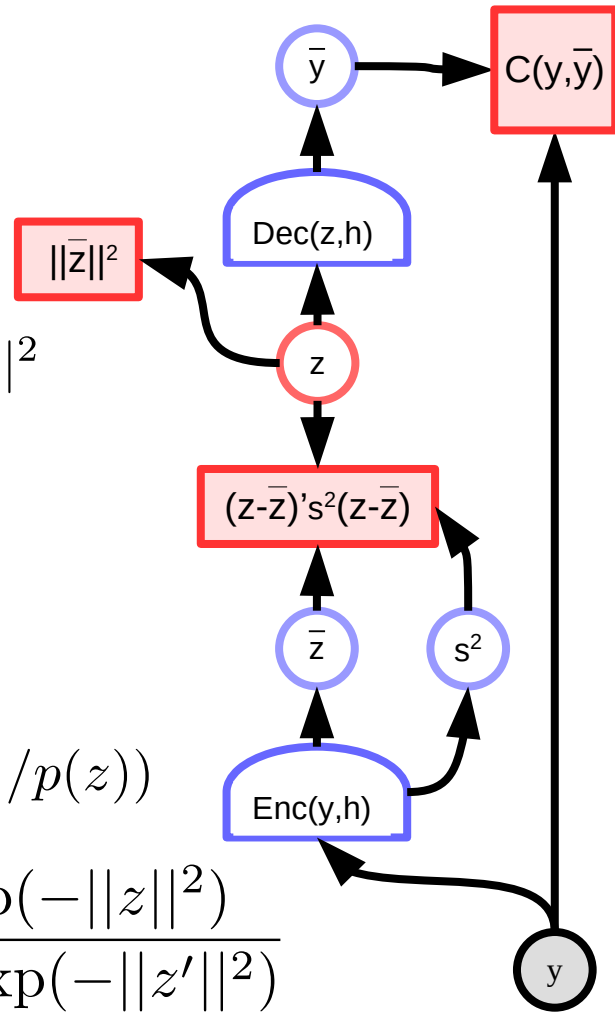
$$Q_{w_e}(y, z) = (z - \text{Enc}_{w_e}(y))^T s^2 (z - \text{Enc}_{w_e}(y)) + \gamma \|\text{Enc}_{w_e}(y)\|^2$$

$$q_{w_e}(z|y) = \frac{e^{-\beta Q_{w_e}(y, z)}}{\int_{z'} e^{-\beta Q_{w_e}(y, z')}} \quad \text{(Gaussian)}$$

$$\tilde{F}_w(y) = \int_z q_{w_e}(z|y) E_{w_d}(y, z) + \frac{1}{\beta} \int_z q_{w_e}(z|y) \log(q_{w_e}(z|y)/p(z))$$


- **Loss** $\mathcal{L}(y, w) = \tilde{F}_w(y)$

$$p(z) = \frac{\exp(-\|z\|^2)}{\int_{z'} \exp(-\|z'\|^2)}$$




Variational Auto-Encoder

► Code vectors for training samples




Variational Auto-Encoder

- ▶ **Code vectors for training sample with Gaussian noise**
 - ▶ Some fuzzy balls overlap, causing bad reconstructions




Variational Auto-Encoder

- ▶ The code vectors want to move away from each other to minimize reconstruction error
- ▶ But that does nothing for us




Variational Auto-Encoder

- ▶ **Attach the balls to the center with a spring, so they don't fly away**
 - ▶ Minimize the square distances of the balls to the origin
- ▶ **Center the balls around the origin**
 - ▶ Make the center of mass zero
- ▶ **Make the sizes of the balls close to 1 in each dimension**
 - ▶ Through a so-called KL term



Backprop as Lagrangian Optimization

Optimization under constraints



Reformulating Deep Learning

► Loss

$$L(x, y, w) = C(z_K, y) \text{ such that } z_{k+1} = g_k(z_k, w_k), \quad z_0 = x$$

► Lagrangian for optimization under constraints

$$L(x, y, z, \lambda, w) = C(z_K, y) + \sum_{k=0}^{K-1} \lambda_{k+1}^T (z_{k+1} - g_k(z_k, w_k))$$

► Optimality conditions:

$$\frac{\partial L(x, y, z, \lambda, w)}{\partial z_k} = 0, \quad \frac{\partial L(x, y, z, \lambda, w)}{\partial \lambda_{k+1}} = 0, \quad \frac{\partial L(x, y, z, \lambda, w)}{\partial w_k} = 0$$

Reformulating Deep Learning

$$\frac{\partial L(x, y, z, \lambda, w)}{\partial \lambda_{k+1}} = z_{k+1} - g_k(z_k, w_k) = 0 \implies z_{k+1} = g_k(z_k, w_k)$$

$$\frac{\partial L(x, y, z, \lambda, w)}{\partial z_k} = \lambda_k^T - \lambda_{k+1}^T \frac{\partial g_k(z_k, w_k)}{\partial z_k} = 0 \implies$$

► **Backprop!**

► Lambda is the gradient

$$\lambda_k = \frac{\partial g_{k-1}(z_{k-1}, w_{k-1})}{\partial z_k}^T \lambda_{k+1}$$

$$\frac{\partial L(x, y, z, \lambda, w)}{\partial w_k} = \lambda_{k+1}^T \frac{\partial g_k(z_k, w_k)}{\partial w_k}$$