

7. Advanced Pretraining and Finetuning Techniques

Efficient transformers

Compute cost of transformers

Q, K, V projections:

$$n \times d_e \xrightarrow{\text{linear}} n \times d \quad O(n \times d_e \times d)$$

Scaled dot-product attention:

$$(n \times d)(d \times n) \xrightarrow{\text{matmul}} n \times n \quad O(d \times n^2) \text{ — slow}$$

Feed-forward layer (GPT-2):

$$n \times d \xrightarrow{\text{linear+ReLU}} n \times d_h \xrightarrow{\text{linear+ReLU}} n \times d \quad O(n \times d \times d_h)$$

Improve efficiency of transformers:

- Quantization (training and inference)
- Weight sharing (training)
- Sparsely-activated models (training and inference)
- Pruning (inference)
- Distillation (inference)

Specifically, improve efficiency of self-attention (reduce the $O(n^2)$ time and memory cost):

- Sparsify the attention matrix
 - Deterministic mask:
 - Blockwise self-attention [Qiu et al., 2020], Longformer [Beltagy et al., 2020]: attention within a local window
 - Data-dependent mask
 - Reformer [Kitaev et al., 2020]: attention within an adaptive local window
- Compress the key-value memory
 - Low-rank prediction (self-attention is low rank)
 - Linformer [Wang et al., 2020]: compute self-attention in a lower dimension
 - Attention-based projection
 - Perceiver [Jaegle et al., 2021]: use latent states to compress the KV memory

Efficient finetuning

- Finetune a subset of parameters

- Freezing the first X layers [[Lee et al., 2019](#)]
- BitFit [[Ben-Zaken et al., 2022](#)]: only finetune the bias term (0.1% of the parameters)
- Adapt the frozen pretrained model
 - Adapter [[Houlsby et al., 2019](#)]: insert small networks to the pretrained model
 - LoRA [[Hu et al., 2021](#)]: add low-rank matrices as additional parameters