

11. Aligning Language Models (Basics)

Overview

What is alignment

- Prompting converts a task to a native LM task, but model performance is sensitive to prompts
- Goal: make human-AI communication natural and efficient
- So that we can just ask the model to do any task

Capability vs alignment

- **Capability:** What things is the model able to do?
- **Alignment:** What things does the model choose to do?
 - Align with human values
 - Provide truthful information and express uncertainty
 - Be careful with potentially harmful information
 - Clarify user intentions and preferences

Challenges in alignment

- **Implicit rules:** [not articulated](#) but assumed in human interaction
 - e.g. Explicit task: answer questions on topic X
Implicit rules: Don't make up stuff. Don't use toxic language. Don't give information that's potentially harmful.
 - The implicit rules may be [context dependent](#):
 - Translation: what if the source text is toxic?
- **Oversight:** provide supervision on alignment
 - One obvious way to align models is to train them on supervised data (later)
 - But how can we supervise models on tasks that beyond human capabilities?
- **Diversity:** whose values should the model be aligned with?
 - Different (cultural/ethnic/gender/religious/etc.) groups agree with different answers to the same question

Approaches to alignment

- **Prompting:** [ask](#) the model to behave according to human values
- **Finetuning / Supervised learning:** [show](#) the model the right response in various context
- **Reinforcement learning:** [reward](#) / [punish](#) the model when its behavior is aligned / unaligned with humans

Prompting

Prompting the model to answer questions truthfully

Prompts can be overwritten — ask it to ignore previous prompts

Summary

Prompt engineering: instruct the model to behave in a certain way

Pros:

- Easy to do—anyone can play around with it
- Efficient—no parameter updates
- First thing to try

Cons:

- Unprincipled—no idea why it works or doesn't work
- Unreliable—performance can have high variance
- Unsafe—easy to bypass

Supervised finetuning

- How do we teach the model the right behavior?
- Going back to supervised learning: **demonstrate** the right behavior
 - Input: user prompt (task specification)
 - Output: (aligned) response
- **Key challenge:** data collection

How to get the prompts and responses?

What kind of data do we need?

Idea 1: use existing NLP benchmarks


But this is not what we ask ChatGPT to do! Distribution shift.

- **Problem:** Gap between training and test data
- Straightforward **solution:** collect training data that is similar to test data

How do we know what test data is like?

- Get some **pilot data**

which requires a working-ish model first!



Tricky cases

- Recall that we want the model to [infer user intention](#)
- But also to make the right decisions that [align with human values](#)
- So it's important to include examples that involve alignment decisions
- Open question: how to handle [trade-off between helpfulness and harmfulness](#)?
e.g., user may request to generate toxic sentences for data augmentation

Summary

Supervised finetuning: train the model to respond in an aligned way on human-annotated prompt-response data

Pros:

- Relatively reliable—generalize to unseen data
- User friendly—doesn't require extensive prompt engineering
- Simple training pipeline—standard finetuning

Cons:

- Need a warm start—pilot data to decide what data to collect
- Expensive—data needs to cover many use cases
- Compute—need to update very large models

Reinforcement learning

Motivation:

- Demonstrations are expensive to obtain—can we learn from weaker signals?
- For many tasks, humans (and animals) only get signal on whether they succeeded or not

Goal: learning from experience by maximizing the expected reward

At each time step t , an agent

- is in a **state** $s_t \in \mathcal{S}$ (\mathcal{S} is the state space)

- takes an **action** $a_t \in A$ (A is the action space)
- transitions to the next state s_{t+1} according to a **transition function** $p(\cdot | s_t, a_t)$
- obtains a **reward** $r(s_t, a_t)$ according to the **reward function** $r : S \times A \rightarrow \mathbb{R}$

The agent uses a **policy** π to decide which actions to take in a state:

- Deterministic: $\pi(s) = a$
- Stochastic: $\pi(a|s) = \mathbb{P}(A = a | S = s)$

A policy π_θ defines a distribution $p_\theta(\tau)$ over **trajectories** $\tau = (a_1, s_1, \dots, a_T, s_T)$.

The agent's **objective** is to learn a policy π_θ (parametrized by θ) that maximizes the **expected return**:

maximize $\mathbb{E}_{\gamma \sim p_\theta(\gamma)} [\sum_{t=1}^T r(s_t, a_t)]$

Key steps:

- **Trial**: run policy to generate trajectories
- **Error**: estimate expected return
- **Learn**: improve the policy

Challenges:

- Trials could be expensive (e.g., healthcare, education)
- Reward signal could be expensive and sparse (e.g., expert feedback)
- May need many samples to learn a good policy

Policy gradient algorithms

While not converged

1. Sample trajectories from the current policy
2. Estimate return for each trajectories based on observed rewards
3. Take a gradient step on the expected return (w.r.t. the policy)

Notation: let $r(\tau) = \sum_{t=1}^T r(a_t, s_t)$ be the return.

Our objective: $J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} [r(\tau)] = \sum_{\tau} p_\theta(\tau) r(\tau)$

$$\begin{aligned} \nabla_\theta J(\theta) &= \nabla_\theta \sum_{\tau} p_\theta(\tau) r(\tau) \\ &= \sum_{\tau} \nabla_\theta p_\theta(\tau) r(\tau) \\ &= \sum_{\tau} p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) r(\tau) \\ &= \mathbb{E}_{\tau \sim p_\theta(\tau)} [\nabla_\theta \log p_\theta(\tau) r(\tau)] \end{aligned}$$

log derivative trick

$$\begin{aligned} &p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) \\ &= p_\theta(\tau) \frac{\nabla_\theta p_\theta(\tau)}{p_\theta(\tau)} \\ &= \nabla_\theta p_\theta(\tau) \end{aligned}$$

Good news: the gradient is now inside the expectation

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p_{\theta}(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) r(\tau)] \quad \text{average gradient of sampled trajectory}$$

But what is $p_{\theta}(\tau)$?

$$p_{\theta}(\tau) = p_{\theta}(a_1, s_1, \dots, a_T, s_T) = p(s_1) \prod_{t=1}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

$$\log p_{\theta}(\tau) = \log p(s_1) + \sum_{t=1}^T \log \pi_{\theta}(a_t | s_t) + \log p(s_{t+1} | s_t, a_t)$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[\left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

Putting everything together

REINFORCE algorithm:

1. Sample N trajectories τ^1, \dots, τ^N from π_{θ}
2. Estimate the gradient:

$$\nabla_{\theta} J(\theta) \approx \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

3. Update the policy with gradient ascent: $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$
4. Go back to 1

How is all this related to LLMs?

Think of tokens as actions:

- Action space: vocabulary $a_t = x_t \in \mathcal{V}$
- State space: history / prefix $s_t = (x_1, \dots, x_{t-1})$
- Policy: a language model $p_{\theta}(x_t | x_{<t})$
- Trajectory: a sentence / generation x_1, \dots, x_T

REINFORCE algorithm on text:

1. Sample N generations from the language model p_θ
2. Estimate the gradient: $\nabla_\theta J(\theta) \approx \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_\theta \log p_\theta(x_t^i | x_{<t}^i) \right) r(x_{1:T})$
3. Update the policy with gradient ascent: $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
4. Go back to 1

What is the algorithm doing?

If $r(x_{1:T})$ is **positive**, take a gradient step to **increase** $p_\theta(x_{1:T})$.

If $r(x_{1:T})$ is **negative**, take a gradient step to **decrease** $p_\theta(x_{1:T})$.

Supervised learning on model generations weighted by rewards

How to get the reward? Next lecture!

Summary

Reinforcement learning: align the model by giving it feedback on whether an output is good or bad

Pros:

- Cost-efficient—humans only need to provide judgments/rewards
- General—can be used to model all kinds of human preferences

Cons:

- Complex pipeline—RL algorithms need more engineering
- Reward hacking—models are good at finding ways to "cheat"

Generating polite and authoritative nonsense

- Human judgments on some subjects are inherently diverse