

2. Process Management

Life of a C program

Preprocessor

The preprocessor expands **directives** such as `#include` , `#define` , `#ifdef` ...

Compiler

The compiler takes the expanded C code, checks the syntax, and generates... Assembly code (gcc); LLVM IR (clang).

In the meantime, it also **optimizes** the code. (gcc -S -O1/O2/O3 program.c)

Assembler

The assembler converts the generated assembly code to an **object file**.

The object file contains **machine code**, but **isn't yet executable**.

Linker

The linker combines **object files** and **libraries** to produce the **executable file**.

A library is just a collection of functions and variables.

Static & dynamic linking: the final program embeds all library functions vs. load dynamic libraries when the program runs.

Processes

The process is the most central concept in an operating system.

- It's an abstraction of a **running program**.
- It attaches to all the **memory** that is allocated for the process.
- It associates with all the **files** opened by the process.
- It contains **accounting information** such as its owner, running time, memory usage...

Let's start with some system calls.

getpid() and getppid()

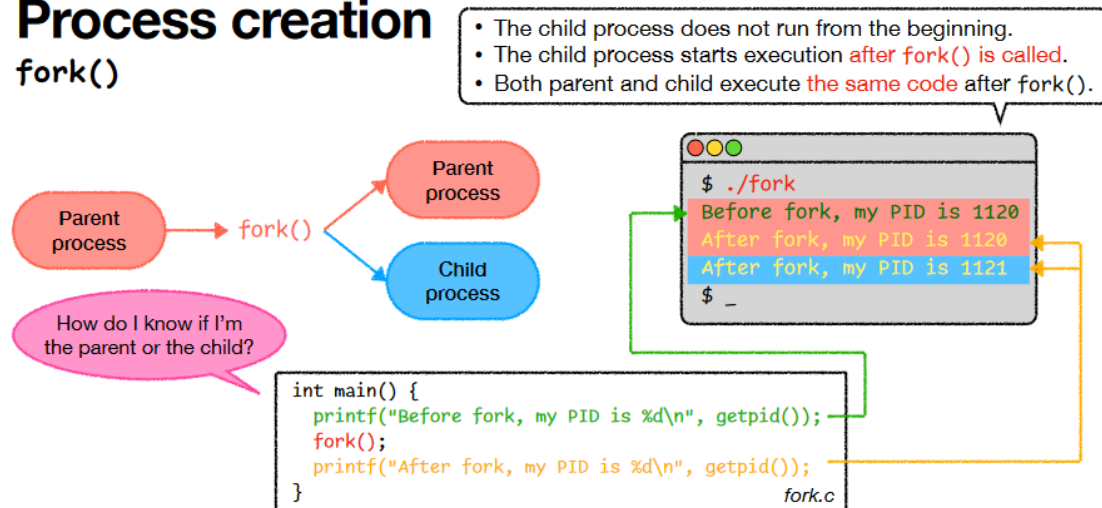
The OS gives each process a unique identification number, the **Process ID (PID)**.

- **getpid()** returns the PID of the calling process.
- **getppid()** returns the PID of the **parent** of the calling process.

fork()

Process creation

fork()



parent process → parent process + child process, executing the same code

To know in which process — `fork()`'s return value differs for the parent and the child:

- In the **parent**, `fork()` returns **the PID of the child process** `> 0`
- In the **child**, `fork()` returns `0`

The child inherits the parent's... Program code; Memory. Differs from the parent in... Return value of `fork()`; Process ID; Parent; Running time.

fork buffer: `printf()` invokes the `write()` system call; **buffer** combines several `printf()` for one `write()` system call, to reduce the number of system calls; at `fork()`, the child inherits the buffer.

Unbuffered: invoke `write()` immediately. Line-buffered: invoke `write()` when a newline character. Fully-buffered: invoke `write()` when the buffer becomes full or before the process terminates.

On a single-processor system, only one process can be executed at one time. After `fork()`, does the parent or the child run first? — We don't know. Decided by process scheduling.

execve() and the exec*() family of functions

When `execve()` is called, the process **replaces** the code that it's executing with the new program and **never returns to the original code** — nothing after `execve()` will be executed.

`exec`**l**, `exec`**p**, `exec`**e**, `exec`**v**, `exec`**vp**, `exec`**ve**.

Path name or file name

- Default — path name, e.g., `"/bin/ls"`.
- **p** — file name, e.g., `"ls"`.

Argument list or array

- **l** — list, e.g., `exec("bin/ls", "bin/ls", "-a", "-l", NULL);`
- **v** — array, e.g., `execv("bin/ls", argv);`

Environment variables

- Default — inherit the current environment.

- `e` — specify a new environment array.

Environment variables are a set of strings maintained by the `shell`.

`PATH=/user/bin; SHELL=/bin/bash; USER=cj2133; HOME=...; PWD=...`

Process creation and execution & `system()`

Can we just run a program and be able to come back?

Yes, there is a convenient library function (not a system call): `system()`

It is implemented using `fork()` and `exec()` to call `/bin/sh -c command`

```
int my_system(const char *command) {
    if (fork() == 0) { // in child
        execl("/bin/sh", "/bin/sh", "-c", command, NULL);
        exit(-1);      // exits if execl() fails, returns -1 to parent
    }
    wait(NULL);        // waits until child terminates
    return 0;
}
```

`wait()` suspends execution of the calling process until one of its children terminates.

`waitpid()` waits for a particular child; waits for a stopped/resumed child.

Summary

A process is created by cloning.

- `fork()` is the system call that clones processes.
- Cloning is copying — the new process inherits many things but not all.

Program execution is not trivial.

- A `process` is the entity that hosts a program and runs it.
- A process can run more than one program.
- The `exec*()` system call family replaces the program that a process is running.

Processes in the kernel

Kernel-space memory vs. user-space memory

Processes in the kernel — Task list ("doubly-linked list"), "node" — process control block, contains everything related to a process.

How to redirect stdout to a file? — `File descriptors`, print to the file instead of the screen.

Process execution

Going back and forth...

A process switches its execution from `user mode` to `kernel mode` by invoking `system calls`.

When the **system call** finishes, the execution switches from **kernel mode** back to **user mode**.

Handling system calls

Example: the `getpid()` system call

The CPU is running a process in **user mode**.

It wants to invoke the `getpid()` **system call**. Each system call has a unique **syscall number**.

The process puts the **syscall number** of `getpid()` in a specific CPU register.

Then it executes a **TRAP** instruction to switch from **user mode** to **kernel mode**.

The kernel starts execution at the **syscall dispatcher**. It examines the **syscall number**, looks up the syscall table, and invokes the corresponding **syscall handler**.

The `sys_getpid()` handler reads the **Process ID** of the calling process from `task_struct`.

Then it executes a **RETURN-FROM-TRAP** instruction to switch from **kernel mode** to **user mode**.

Execution of a process: user time vs. system time

Process-management syscalls in the kernel

`fork()`

`execve()`

`wait()` & `waitpid()` & `exit()` → child enters the zombie state

Signals

Signals are software interrupts. It's a form of **inter-process communication (IPC)**.

A process can send a signal to another process.

When a signal arrives, the OS interrupts the target process's normal control flow and executes the **signal handler**.

How are signals generated?

From the user space (keyboard, commands, `kill()`...) / the kernel or hardware (`SIGCHLD`...)

The **kill** command sends **SIGTERM** to the target process by default. The default signal handler of **SIGTERM** is to **terminate the process**.

Foreground and background jobs run concurrently. **fg**: resumes a job in the foreground (send a **SIGCONT** signal and then wait for the child).

Just like the **kill** command, the **kill()** syscall sends a signal to a process.

The **raise()** library function sends a signal to itself (by calling `kill()` with its pid).

The **signal()** system call changes the **signal handler**.

The **pause()** system call puts the process to **block (sleep)** until the delivery of a signal handled by the process, or a signal terminating the process.

Process organization

Booting the computer

BIOS (Basic input/output system) is firmware (i.e., software providing low-level control for hardware).
Locates the **boot device** → boot code determines the bootable partition → executes boot loader → starts the OS kernel → creates **the first process (init/systemd)**

Orphans: If the shell terminates, the program becomes an orphan. The **init** process automatically adopts all orphaned processes (becomes their new parent).

Processes in Linux are organized as **a single tree** (instead of a forest).

Reparenting allows processes to run **without a parent shell**.

Process scheduling

Multiprogramming: Computers often do several things concurrently, even if it has only one CPU.

Multitasking: The CPU switches from process to process quickly, running each for a few ms.

Scheduling : The OS needs to **choose** which process to run next. The part of the OS that makes the choice is called the **scheduler**.

Process states

Ready: **runnable but temporarily stopped** to let another process run.

Running: the process is actually **using the CPU** at that instant.

Blocked: **unable to run** until some **external event** happens. interruptible/uninterruptible.

Zombie: the process **exits** and its **parent has not yet waited** for it.

The **context** of a process consists of its user-space memory; register values.

The **scheduler** decides which process to run next. If the next process is different from the current one, the OS performs a **context switch** (saving and restoring registers; switching memory maps...).

Most processes' execution **alternates between CPU execution and I/O wait**.

CPU-bound processes spend most of their running time on the CPU. user time > sys time.

I/O-bound processes spend most of their running time on I/O. user time < sys time.

Nonpreemptive scheduling: a process keeps running until... it starts waiting for I/O; or it voluntarily relinquishes the CPU.

Preemptive scheduling: a process can run for a particular period of time, then suspended and another process may run.

What's a good scheduling — general goals

- Fairness: each process should have a fair share of the CPU.

- Policy enforcement: the system's policies should be carried out.
- Balance: all parts of the system should be kept busy all the time.

Batch systems: process jobs in batches, goals: high throughput, short turnaround time

Interactive systems: have interactive users, **preemption** is essential, goal: short response time

Real-time systems: guarantee response **within specified time constraints**, goal: meet deadlines

Scheduling algorithms

Metrics

Wait time — the duration that the job is in the system but **not running**

Turnaround time — the duration from when the job **arrives** in the system to it **completes**. $T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$

Response time (in **interactive systems**) — the duration from when the job **arrives** in the system to **the first time it is scheduled**. $T_{\text{response}} = T_{\text{first run}} - T_{\text{arrival}}$

First-come, first-served (FCFS), a.k.a. first-in, first-out (FIFO), is a **nonpreemptive** scheduling algorithm in **batch systems**.

Shortest job first (SJF) has both **nonpreemptive** and **preemptive (shortest remaining time)** versions.

Round-robin (RR) is a **preemptive** scheduling algorithm in **interactive systems**.

Each job is assigned a **time slice** (a.k.a. **quantum**), the amount of time the job is allowed to run. Jobs are running one by one in the **run queue**.

RR has worse CPU efficiency than SJF. However, jobs on a RR scheduler are more **responsive**. This is an inherent **trade-off** between performance (SJF) and fairness (RR). RR needs more context switching, which is relatively slow.

Priority scheduling: Each job is assigned a **priority**. The scheduler always chooses the job with the **highest priority** to run. Priorities can be **static** or **dynamic**.

Static priority scheduling limitations: low-priority jobs may starve to death.

Dynamic priority scheduling: no standard way.

Multilevel feedback queue (MLFQ): It's a kind of dynamic priority scheduling, but **each priority has its own policy**. MLFQ **observes how jobs behave over time, and prioritize them accordingly**.

All jobs start running at the highest priority. When a job uses up its time slice, its priority is reduced by 1. If a job gives up the CPU before the time slice is up, it stays at the same priority.

Limitation if a program changes its behavior over time → new rule: after some time period, move all jobs to the highest priority (priority boost).

Interprocess communication

Processes often need to **communicate** with one another.

IPC: to share information, to reuse software, to speedup computation (e.g. MapReduce).

Case study: the `pipe()` system call returns **two file descriptors**: `pipefd[0]` and `pipefd[1]`. This is called a **producer-consumer model**.

The **kernel** provides a form of **synchronization**. However, for **shared memory**, it's up to the **processes** to coordinate.

Race condition: race to access the shared resource. The results depend on the **timing** of the execution, i.e., the particular **order** in which the **shared resource** is accessed.

To avoid race conditions, we need **mutual exclusion**: If one process is accessing a shared resource, the other processes must be excluded.

A **critical section** (a.k.a. critical region) is a piece of code that accesses a shared resource. It should be as tight as possible; it may access multiple shared resources.

Requirements:

- No two processes may be simultaneously inside their critical sections — mutual exclusion
- No assumptions may be made about the speeds or the number of CPUs
- No process outside critical section may block others → all process can make progress
- No process should have to wait forever to enter its critical section → **bounded waiting**

Spinlocks are built upon CPU instructions that guarantee **atomic operation**.

A semaphore is an object with a **non-negative** integer value.

The value **must be initialized** before being used.

There are two operations: **down()** and **up()**

```
void down(semaphore *sem) {
    // atomic operation
    if (*sem > 0) {
        *sem = *sem - 1;
    } else {
        block on sem;
    }
}

void up(semaphore *sem) {
    // atomic operation
    if (some process is blocked on sem) {
        let one such process proceed
    } else {
        *sem = *sem + 1;
    }
}
```

Classical IPC problems

The producer-consumer problem

- It models access to a bounded buffer.

The dining philosopher problem

- It models processes competing for exclusive access to a limited number of resources (e.g., I/O devices).

The readers and writers problem

- It models access to a database.

The sleeping barber problem

- It models a queueing situation

The producer-consumer problem (the bounded-buffer problem)

Two processes **share a fixed-size buffer**.

The **producer** **inserts** data to the **tail** of the buffer.

The **consumer** **removes** data from the **head** of the buffer.

Requirements:

1. When the producer wants to insert an item into the buffer, but **the buffer is already full...**
 - The producer should **block**.
 - The consumer should **wake up** the producer **after it has consumed an item**.
2. When the consumer wants to remove an item from the buffer, but **the buffer is empty...**
 - The consumer should **block**.
 - The producer should **wake up** the consumer **after it has produced an item**.
3. **Mutual exclusion**: no two processes can access the shared buffer at the same time.

Implementation using semaphores

```
semaphore mutex = 1;           // controls access to critical section
semaphore empty = BUFFER_SIZE; // counts empty buffer slots
semaphore filled = 0;          // counts filled buffer slots
```

Why do we need three semaphores?

```
1 void producer() {
2   int item;
3
4   for (;;) {
5     item = produce_item();
6     down(&empty); // decrement empty count
7     down(&mutex); // enter critical section
8     insert(item); // put new item in buffer
9     up(&mutex);   // exit critical section
10    up(&filled);  // increment filled count
11  }
12 }
```

```
1 void consumer() {
2   int item;
3
4   for (;;) {
5     down(&filled); // decrement filled count
6     down(&mutex);  // enter critical section
7     item = remove(); // take item from buffer
8     up(&mutex);    // exit critical section
9     up(&empty);    // increment empty count
10    consume_item(item);
11  }
12 }
```


mutex guarantees **mutual exclusion**.

empty and **filled** are for **synchronization**.

empty represents the number of **empty slots**.

filled represents the number of **filled slots**.

Invariant: **empty** + **filled** = buffer size


If we swapped lines 6 & 7 in the producer, the scenario is called a **deadlock**.

It happens when there is a **circular wait**...

- The producer is waiting for the consumer to up() the **empty** semaphore.
- The consumer is waiting for the producer to up() the **mutex** semaphore.

The dining philosopher problem

The dining philosopher problem




5 philosophers around a table;
5 instant ramen noodles;
5 chopsticks in between.

A philosopher does only two things in their entire life:
eating and **thinking**.

In order to eat, a philosopher needs **exactly two chopsticks**.

How to design a protocol so that they can enjoy their life?



For each philosopher i ,

- **left chopstick** = i
- **right chopstick** = $(i + 1) \% N$
- **LEFT neighbor** = $(i + N - 1) \% N$
- **RIGHT neighbor** = $(i + 1) \% N$

States:

- **THINKING**
- **HUNGRY** ← Trying to get chopsticks.
- **EATING** ← Critical section.

Requirements:

1. Mutual exclusion
2. Synchronization: the solution should avoid any potential **deadlock** or **starvation** scenarios.

Potential issues:


- Modeling each chopstick as a semaphore: we are afraid to down() a chopstick, as that may lead to a deadlock.
- However, our real issue is that philosophers need to eat, not chopsticks. We should guarantee: when a philosopher is eating, their left and right neighbors cannot eat.

<pre>int state[N]; // THINKING, HUNGRY, or EATING semaphore mutex = 1; semaphore sem[N]; // what are the initial values?</pre>	<pre>void test(i) { if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) { state[i] = EATING; up(&sem[i]); // take both chopsticks } }</pre>
<pre>void take_chopsticks(int i) { down(&mutex); state[i] = HUNGRY; test(i); // try to take both chopsticks up(&mutex); down(&sem[i]); // block if cannot take chopsticks }</pre>	<pre>void philosopher(int i) { for (;;) { think(); take_chopsticks(i); // section entry // block until I take both chopsticks eat(); // critical section put_chopsticks(i); // section exit } }</pre>
<pre>void put_chopsticks(int i) { down(&mutex); state[i] = THINKING; test(LEFT); // see if left neighbor can now eat test(RIGHT); // see if right neighbor can now eat up(&mutex); }</pre>	

The readers and writers problem

The readers and writers problem

Accessing a database



Multiple processes are allowed to **read** the database at the same time.

If a process is **writing** the database, no other processes may have access to the database.

How to program the **readers** and **writers**?

Requirements:

1. Mutual exclusion: the database is a shared resource.
2. Synchronization
 - When a reader is reading, other readers are allowed to read the database.
 - When a reader is reading, no writers are allowed to write the database.
 - When a writer is writing, no readers or writers are allowed to access the database.

3. Concurrency: concurrent access from multiple readers should be allowed.

```

semaphore db = 1;    // controls access to the database
semaphore mutex = 1; // controls access to "reader_count"
int reader_count = 0; // # of processes reading or wanting to read

void reader() {
    for (;;) {
        down(&mutex);
        if (++reader_count == 1)
            down(&db); // the first reader locks db
        up(&mutex);
        data = read_database(); // critical section
        down(&mutex);
        if (--reader_count == 0)
            up(&db); // the last reader unlocks db
        up(&mutex);
        consume_data(data);
    }
}

void writer() {
    for (;;) {
        data = produce_data();
        down(&db); // locks db
        write_database(data); // critical section
        up(&db); // unlocks db
    }
}

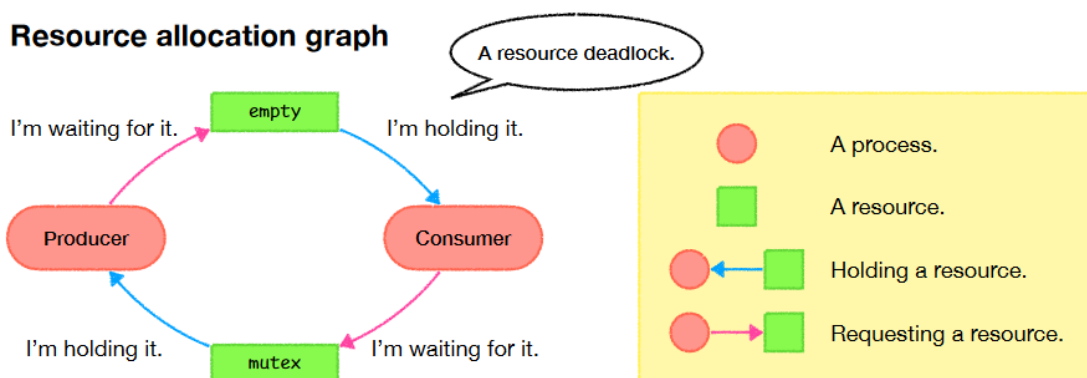
```

This solution meets all our requirements.
 However, it gives readers a higher priority than writers.
 As long as there is a steady supply of readers, writers will suffer from starvation.

Deadlocks

Remember the deadlock in the producer-consumer problem?

Resource allocation graph



Conditions for resource deadlocks:

Mutual exclusion

- Each resource is either available or currently assigned to **exactly one process**.

Hold and wait

- Processes currently holding resources that were granted earlier **can request new resources**.

No preemption

- Resources previously granted **cannot be forcibly taken away** from a process.
- They must be explicitly released by the process holding them.

Circular wait

- A **cycle** of 2+ processes, each of which is waiting for a resource held by the next member of the cycle.

Preventing deadlocks

The most adopted approach is the ostrich algorithm — stick your head in the sand and pretend there is no problem.

Attack #1: no mutual exclusion?

Method #1: make data **read only**; so processes can use the data concurrently.

Method #2: implement **lock-free** data structures.

- test-and-set; compare-and-swap (CAS)

Method #3: **read-copy-update (RCU)**.

- It allows a writer to update the data structure while other processes are still using it.
- A reader would see and traverse either the old version or the new version.

Attack #2: no hold-and-wait?

Method #1: **request all resources** at once before starting execution.

- Drawback: a process may not know what resources they will need; resources held for longer than needed, thereby decreasing concurrency.

Method #2: **release all resources** before requesting a new one.

Attack #3: no “no preemption”?

Technically, we can't forcibly take away a resource that a process already holds. In practice, a process can “**preempt**” **their own ownership** in a graceful way. We used this approach in the diningphilosopher problem. It's called **trylock**.

Attack #4: no circular wait?

It's probably the most practical and frequently used approach.

All the resources are given a **total order** (i.e., global numbering). A process can request only resources **higher** than what it's already holding.

Detection and recovery

Allow deadlocks to occur (occasionally); Try to detect when this happens; take action to recover

Detecting deadlocks

A deadlock detector runs periodically and build the **resource allocation graph**; **find a cycle**.

Threads: lightweight processes

A thread is an **execution entity within a process**.

A **multithreaded process** can have more than one execution in it

- All threads **share the same code**.
- A new thread starts with a specific **thread function**.
- The thread function can invoke other functions and system calls.
- However, the thread function **does not return to its caller**.

All threads **share the same global variables** and **dynamically allocated memory**.

Each thread has its own stack for **local variables**. We can still access another thread's stack if we know the memory address.

It allows **multitasking** within a process:

- One thread waits for the user input;
- Another thread performs computation.

→ Better **performance** and **responsiveness**.

Threads are **easier to create and destroy** than processes (10-100× faster).

Threads share the **same address space**; so sharing data is easy.

Thread models

Many-to-one model

- Implement threads in **user space**.
- Cons: When a **blocking system call** is invoked, all threads will be blocked; **Page faults** (discussed later) in a thread will block the entire process; No preemption of threads due to the absence of clock interrupts.

One-to-one model

- Implement threads in **the kernel**.
- Pros: When a thread blocks, the kernel can **schedule another thread**.
- Cons: Creating and destroying threads are more expensive; trick for forks & signals

Many-to-many model

- **Hybrid implementations**.
- great, only complex to implement

Thread programming

POSIX threads (man 7 pthreads)

Description	Process	Thread
Process/thread identification.	pid_t	pthread_t
Get process/thread ID.	getpid()	pthread_self()
Create a new process/thread.	fork()	pthread_create()
Terminate the calling process/thread.	exit()	pthread_exit()
Wait for a specific process/thread to exit.	waitpid()	pthread_join()
Send a signal to a process/thread.	kill()	pthread_kill()
Release the CPU to let another process/thread run.	sched_yield()	pthread_yield()

Thread programming

Mutual exclusion

A **mutex object** of type `pthread_mutex_t` is similar to a **binary semaphore**.

Description	pthread mutex	Equivalent semaphore
Initialize a mutex.	<code>pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;</code>	<code>semaphore mutex = 1;</code>
Lock a mutex. (Enter critical section.)	<code>pthread_mutex_lock(&mutex);</code>	<code>down(&mutex);</code>
Unlock a mutex. (Exit critical section.)	<code>pthread_mutex_unlock(&mutex);</code>	<code>up(&mutex);</code>

Note: `pthread_mutex_t` cannot be used as counting semaphores.

Thread programming

Semaphores

A **semaphore object** of type `sem_t` is similar to a **counting semaphore**.

Description	POSIX semaphore	Equivalent CS202 semaphore
Initialize a semaphore.	<code>sem_t s;</code> <code>sem_init(&s, 0, initial_value);</code>	<code>semaphore s = initial_value;</code>
Down a semaphore.	<code>sem_wait(&s);</code>	<code>down(&s);</code>
Up a semaphore.	<code>sem_post(&s);</code>	<code>up(&s);</code>
Destroy a semaphore.	<code>sem_destroy(&s);</code>	

Thread programming

Condition variables

Although **semaphores** provide a convenient and effective synchronization mechanism, **using them incorrectly can cause hard-to-detect timing errors**.

The pthread library provides another synchronization mechanism: **condition variables** (of type `pthread_cond_t`).

- Later, we'll see that **semaphores** can be easily implemented using **mutexes** and **condition variables**.

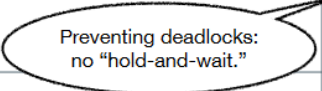
Condition variables allow a thread to **wait for a condition to become true**.

Example: the producer-consumer problem.

- If the buffer is full, the **producer** needs to **block** and **be awakened when the buffer is not full**.
- If the buffer is empty, the **consumer** needs to **block** and **be awakened when the buffer is not empty**.

Thread programming

Condition variables

Usage	Description
<code>pthread_cond_t cond = PTHREAD_COND_INITIALIZER;</code>	Initialize a condition variable.
<code>pthread_cond_wait(&cond, &mutex);</code> 	Wait on a condition. The calling thread must have locked the mutex . <ul style="list-style-type: none">When the thread is about to block, mutex will be unlocked automatically and atomically.When the thread is unblocked, mutex will be locked again, automatically and atomically.
<code>pthread_cond_signal(&cond);</code>	Signal a condition. <ul style="list-style-type: none">If some threads are blocked on cond, at least one thread will be unblocked.If no thread is blocked on cond, the signal is lost.

A thread pool is a **design pattern** for achieving concurrency of execution.

It maintains a **pool of worker threads** waiting for tasks to be dispatched.

Some data structures and functions are designed for single-threaded execution.

- Example: `strtok()` uses a static buffer while parsing a string, thus not thread safe.

One way to achieve thread safety is to make the function **reentrant**.