

# 3. File Systems

We will look at file systems (FS) from two viewpoints:

- Static view: what's the **layout** of a file system?
- Dynamic view: what **operations** does a file system support?

An OS can support more than one FS.

A FS can be accessed by more than one OS.

## Files

Why do we need files?

Storing information in memory is good since memory is fast, but memory vanishes after reboot.

Files provide **long-term** information storage. They are persistent.

Files can also be **shared objects** for processes to access concurrently.

Filenames: When a file is created, a **name** must be given.

Pathname vs. Filename: the pathname is **unique** across the entire file system; the filename is **not unique**, but **unique** within the directory.

File attributes: important to the file system. **FS-dependent**, not OS-dependent. **stat()** system call can read file attributes, or check whether a file exists or not without opening that file.

File types

- Regular file: inspect text files using cat; binary files using xxd. The command **file** can distinguish the type of a file, without checking the extension of the filename.
- Directory file
- Link file
- Device file

Distinguishing binary files: magic number

**Directories:** A directory is also a **file**. A directory file is **an array of directory entries**. It records all the files and directories that belongs to it.

**Links:** A link is an **alias** of a file.

- In Windows, we have the **shortcut**.
- In Unix/Linux and macOS, we have the **hard link** and **symbolic link**.

Both the shortcut and the symbolic link share the same concept.

- A link file is created and stores the **pathname** of another file.
- To create a symbolic link: `ln -s filename linkname`

We will come back to the hard link when we discuss the file system implementation.

Device files: Unix/Linux creates some pseudo-devices for convenience.

- /dev/zero: output an endless sequences of zeros.
- /dev/urandom: output an endless sequence of random bytes.
- /dev/null (a.k.a. the black hole): whatever data written to this device will be discarded.

File creation = update of the directory file

File deletion: Removing a file is the reverse of the creation process.

## File system design

A file system is about...

- How the OS stores and locates a **file**.
- How the OS stores and locates a **directory**.

It's also about how the OS manages the storage in an **efficient** and **reliable** way.

- Efficiency is about the speed and the storage utilization.
- Reliability is about “can we always get back what have been saved?”

File system structure:

- Application level (user)
- Logical representation level (partially user and mostly kernel)
- File organization level (kernel)
- Device control level (kernel)

File organization: goal is to store files on the disk, read and write the stored files in an efficient and reliable way

linked list to store the next block at each block → store all the linked lists as a table

**File allocation table (FAT)** approach

- All the information about the next block are **centralized** as a table.
- The entries in the table are stored **contiguously as an array**.
- Each entry corresponds to **one file allocation block** in the storage device.

Drawback: the file allocation table takes a lot of space (thousands of blocks).

In FAT, a block is called a **cluster**. Different versions of FAT support different number of clusters.

$L$  = length of the cluster's address,  $S$  = size of a cluster, then maximum file size =  $2^L \cdot S$ .

### Index node (inode) approach

- The file system table maintains the mapping from the filenames to the index nodes (inodes).
- An inode stores the **addresses of all the data blocks**.
- All the inodes are stored in the **inode table**.

Each inode has a **fixed size**. At the beginning, the addresses of the data block are stored in the **direct block addresses** of the inode. After they have been used up, the **indirect block addresses** will be used, which point to other data blocks. There can be indirect blocks of indirect blocks.

Summary: block-based addressing, one large table (FAT) → many small tables (Inode)

Free-space management: keep track of every free block in the file system

A bitmap of block allocation: a series of boolean values

FAT32: hints of first free cluster & number of free clusters; little space needed but inaccurate & slow

Support directories: a directory is just a list of directory entries

File attributes: stored inside a directory entry (FAT16, FAT32) / inside an inode (ext2, ext3, ext4)

File system information

Boot sector (FAT, NTFS): the first 512 bytes of the FS to store all the FS-specific data

Superblock (ext2, ext3, ext4): a 1024-byte region to store all the FS-specific data

Disk partition: a logical space to host a file system

A smaller file system is more efficient

“formatting” a disk — creating and initializing a file system

## FAT32 in action

### Reading a directory

Step 1: locate the root directory

Step 2: find the directory entry

Step 3: follow cluster number, read and return all directory entries

Step 4: read and return the entire directory entry structure

### Directory entry (32 bytes)

Bytes 0-10: 8+3 characters of filename + extension

Bytes 20-21: high 2 bytes of the first cluster address

Bytes 26-27: low 2 bytes of the first cluster address

Bytes 28-31: file size

### Reading a file

- Step 1: read the first cluster from the directory entry
- Step 2: look up the next cluster in the FAT and read that cluster
- Step 3: repeat the process until an EOF entry is found in the FAT

### Writing a file

- Step 1: for appending, locate the cluster of the end of the file
- Step 2: look up the FSINFO structure to find the next free cluster
- Step 3: allocate new cluster by changing the FATs and FSINFO

### Deleting a file

- Step 1: find the directory entry and the locations of the clusters
- Step 2: set all the next address fields in the FATs of that file to 0
- Step 3: update the FSINFO structure
- Step 4: change the first character of the filename to 0xE5

The file is **not really deleted**! Perform a search in all the free space to find all deleted file contents. Those data persists until the deallocated clusters are **reused**. This is a trade-off between performance (during deletion) and security. There are ways to delete a file securely.

How to recover a deleted file?

**Pull the power plug immediately!** It prevents the target clusters from being **overwritten**.

If the file size is **no more than one cluster**, the recovery can be easily done.

- The first cluster address is still searchable.
- Note that files with the same suffix may also be found.

If the file size is **larger than one cluster**, other clusters' addresses are all gone...

- However, due to next-available search, clusters of a file are likely to be contiguously allocated.
- If not, we'd better have the exact file size and the checksum of the deleted file beforehand so that we can use a brute-force method to recover the file.

### ext2/3/4 in action

The primary file system on Linux is **ext4** (fourth extended filesystem).

default block size = 4KB (configurable from 1KB to 64KB), default block address = 32-bit

The file system is partitioned into **block groups**.

- Each block group has the same internal structure.
- The **group descriptor table (GDT)** stores important information.

Why doing so?

- For **performance** and **reliability**.
- To keep the metadata and the file contents **close together**. Thus, the disk head does not need to travel a long distance.
- The metadata is scattered, so there is **no single point of failure**.

The metadata structure of a block group

**Block bitmap** stores the **block allocation status** of the same block group.

**Inode bitmap** stores the **inode allocation status** of the same block group.

**Inode table** stores the **contents** of all the inodes of the same block group.

- Each inode has a fixed size specified in the superblock.
- The inodes are stored sequentially.

The numbering of inodes starts from 1. The first 10 inodes are reserved for special purposes.

Inode #1 — bad blocks; Inode #2 — root directory; Inode #8 — journal data blocks.

## Link files

Hard links

A **hard link** is a directory entry pointing to an **existing file**.

- **No new file content** is created!

Conceptually, this creates **a file with two filenames**.

- Deleting only one of the directory entries **will not delete the file content!**

The **link count** field in the inode keeps track of how many directory entries are pointing to this file.

- **unlink() system call** to remove a link & link count -= 1; remove file content when link count = 0

Special hard links

- The directory **“.”** is a hard link to **itself**.
- The directory **“..”** is a hard link to its **parent directory**.

When a **regular file** is created, the link count is **always 1**.

When a **directory** is created, the link count is **always 2**.

**Removing a directory** is as simple as **removing the directory entry** and **decrementing its link count**.

- The **rmdir() system call** is used

Symbolic links

A **symbolic link** is a **file**.

- Unlike a hard link, a **new inode is created** for a symbolic link.

Where is the target path stored? It depends on the **length** of the path.

- If the path is fewer than 60 characters, it is stored in the **12 direct block and the 3 indirect block pointers**.
  - $(12 + 3) \times 4 = 60$
- If the path is more than 60 characters, **one data block** is allocated to store the path.

### Writing a file

Step 1: Use a linear search in the inode bitmap to find an **unallocated inode** for the new file.

Step 2: Use a linear search in the block bitmap to find **unallocated data blocks** for the new file.

Step 3: Read the inode of the **root directory**, i.e., inode #2.

Step 4: Following the data block pointers, read the directory entry structure for “/dir1”

Step 5: Read inode #123 and the data blocks. Construct the directory entries.

Step 6: Add a new directory entry with inode #1024 to “/dir1 ”.

### Deleting a file

Step 1: Read the inode and the data blocks of “/dir1 ” and locate the inode of “picture.jpg”.

Step 2: Read the inode #894 and decrement the link count of “picture.jpg ”. Now, the link count has become zero.

Step 3: Deallocate the data block and the inode by setting the corresponding bits in the block bitmap and the inode bitmap.

Step 4: Change the entry length of the previous directory entry of “picture.jpg”.

The change in the entry length aims to **skip the deleted entry**. Such a change produces **holes** among the directory entries. Again, as the same happened in FAT32, the file is actually **not deleted**!

## Linux file system internals

Logic representation level

It's about how the kernel models a file in a logical way.

- How does the kernel interact with the disk?
- How does the kernel interact with the user?

### Kernel-process relationship

— a strong relationship between a **process** and the **opened files**.

The user program is given a **file descriptor** returned from the **open()** system call as an abstract representation of an opened file. It's so abstract that it's just a **number**!

All the opened files are stored inside a structure of type `struct files_struct`.

- By default, every process initially has **three opened files**.
- File descriptors **0** for **stdin** stream, **1** for **stdout** stream, **2** for **stderr** stream

When the process calls the `open()` system call to open a file...

- The **file system module** reads the **metadata** of the file from the disk.
- The opened files will be represented as a structure and **added to the process**.
- A **new file descriptor** will be allocated for the new file.

The structure created contains the information about the **directory entry**. In turn, the directory entry contains the **name** and the **inode** of that file.

The **scope** of a file descriptor is restricted **inside the process**. Different processes may use **different** file descriptors to refer to the same file.

## Virtual file systems

For each opened file, **a set of file system functions** is associated with it.

- They are just **function pointers**. They are called **VFS functions**. They are **FS-specific**.

The VFS functions are invoked by the **FS-related system calls**.

Pros

VFS allows the OS to support **multiple file systems** at the same time.

- You know, the `open()` for ext4 is a lot **different** from the `open()` for FAT32.

It's easy to **support a new file system** in the future.

- All you need to do is **implement the FS functions** based on the **VFS structure**.

Cons

Cannot utilize **special functions** of a specific file system.

## File system performance issues

### Data fragmentation

It's about the **allocated blocks** of a file. A problem arises when the blocks of a file is **randomly scattered** on a disk.

Some basics about the hard disk drive: A hard disk drive has a movable disk arm, which has a disk head for reading and writing data. Data are stored sector-by-sector on the platters using magnetism. The platters are spinning. Therefore, a spinning of the platters can retrieve **a set of continuous sectors**.

If the data blocks are scattered, **the number of disk head movements** will be large. The number of disk head movements determines the **time in accessing the file**.

We need **defragmentation**.

### Cache

To **save disk access time**, we want to keep **frequently accessed** files in the memory at all times. That piece of memory is called a **cache**.

Cache replacement algorithm

## File system consistency and recovery

System crashed; detect & recover the inconsistencies after the operating system comes back

Error #1: Free space is allocated to nothing!

**fsck** can construct a bitmap of allocated blocks using the inode table, thus discover **missing blocks**.

Reset the **free space bitmap**; reset field(s) in the superblock.

Error #2: Allocated space are chained together now. But the space is not associated to any files.

**fsck** can discover missing blocks, and **dangling inodes**, i.e., inodes created without any directory entries referring to

Reset the free space bitmap; reset field(s) in the superblock; delete the dangling inodes.

Error #3: The directory entry is pointing to uninitialized data! The content of the file is wrong!

A **journal** is the log book for the file system. It's kept inside the file system,

A set of file system operations becomes an **atomic transaction**. A **transaction** marks all the changes that **will be done** to the file system. Every transaction is written to the **journal**.

After a crash, the recovery of file system is to **replay the uncommitted transactions on the journal**.

What if a crash happens while the system is writing the journal? During recovery, the transaction is found to be incomplete, and it is not able to be completed.

## RAID

Idea: use **redundancy** to improve both **performance** and **reliability**.

- Use **redundant array of inexpensive disks** as one storage unit.
- Fast: simultaneously read and write disks in the array.
- Reliable: use parity (i.e., XOR) to detect and correct errors.