

4. Memory Management

We'll focus on Intel's memory management.

- Process level (logical level)
- Kernel level (mapping from logical level to physical level)
- Device level (physical level)

Memory Management from a process' perspective

A process is running some **program code**.

The process can change this relationship using the **exec*()** system call family.

The running program code is required to be **in the memory**, since the CPU needs to fetch the instructions from the memory for execution.

Global variables and constants are stored in the **object file**.

The number of global variables and constants are **fixed at compile time**.

When a new program code starts running (started by an **exec*()** system call), the **global variables** and the **constants** are created in the memory.

Constants are **read-only** and should be **protected**.

Local variables are **bound to a function only**.

When a function is **invoked**, the local variables are **created**. When a function **returns**, the local variables are **abandoned**.

Surprisingly, the kernel is not involved during the function calls. The **compiled code** contains all the memory management.

Local variables are organized as a **stack**.

Function **parameter passing** is also done through the stack. A function **refers to the parameters as if they are local variables**.

When a function returns, all the local variables and function parameters are popped out.

The **frame pointer** (a.k.a. base pointer or BP) register points to the stack top as it was when the function was just invoked.

No more space left in the stack — stack overflow. Use dynamically allocated memory (heap).

Allocating memory is done by the **brk()** system call, which grow or shrink the allocated area, or **mmap()** system call, which can map files or devices into memory.

malloc() only **manages** free space returned by the kernel. It **internally** needs memory to manage that, so it asks for more bytes using **brk()** or **mmap()**.

Memory Management from the kernel' perspective

The kernel knows...

- How many processes are in the system.
- How much space each process needs.
- How much memory is in the system.

A portion of the hard disk will be reserved to serve as the memory. We call this the **swap**.

Swapping increases the time in context switching. It causes **external fragmentation** — use defragmentation, a.k.a., **memory compaction**.

Address space: when the CPU wants to read/write a piece of memory, a **memory address** is needed.

The same variable may have **different addresses** at different times (e.g. swapping).

Instead of physical addresses, the compiled code uses **logical addressing**. A logical address can be **translated** to a physical address by the OS kernel or the CPU.

base address (physical address, one for each process) + **offset** (logical address) = target address

Virtual memory

Virtual memory is a memory management technique that solves all these problems...

- External fragmentation.
- Memory growth.
- Memory size limit.

Virtual memory allows...

- The **physical address space** of processes can be **noncontiguous**.
- Processes can run even when they are only **partially** in the memory!

Every process has the **entire** address space.

The **physical memory** is partitioned into fixed-size (e.g. 4KB) blocks called **frames** (page frames).

A process' **logical memory** is also partitioned into blocks of the same size called **pages**.

The virtual memory allows a process' memory to **grow**. The free space inside the process is **so virtual that it's not even pre-allocated**. It's indeed **free space**.

Virtual memory internals

The big problem: the CPU runs in a fetch-decode-execute cycle.

We need to **translate** virtual addresses into physical addresses by **memory management unit (MMU)**, a chip inside the CPU.

Page table: stores the **memory mapping**.

- It tells which pages are in the physical memory.
- It's **stored in the memory** and **used by the MMU**.

The content of the page table depends on the **paging mechanism**.

Demand paging loads a page from the disk to the physical memory **when that page is demanded**.

It's similar to swapping, but the system **will not load the entire process** into the memory, but **only the pages required**. This allows more processes to be hosted in the system.

A **translation lookaside buffer (TLB)** is an address-translation **cache** inside the MMU. It stores the **recent translations** of virtual memory to physical memory.

The page table must be contiguous → too large.

Multilevel page tables: a page table of page tables.

Page replacement

Condition 1: The process requests **a page that does not exist** in the physical memory.

Condition 2: There are **no free frames** available in the physical memory

Step 1: decide the **victim page** in the physical memory

Step 2: **swap out** the victim and write it back to the disk

Step 3: **swap in** the desired frame

The kernel needs to know...

- The current **frame allocation status**.
- **page reference string**: the **order of the pages** that are being referenced by the running process.

Page replacement algorithm: to locate which page is the victim page.

- **First-in first-out (FIFO)** page replacement.
- **Optimal** page replacement: replace the page that **will not be used for the longest period of time** in the future, but impossible since we don't know the future.
- **Least-recently-used (LRU)** page replacement.
 - Give every frame an **age**.
 - If the frame is **just used**, reset the frame's age to **0**.
 - **Other frame's** ages are **incremented by 1**.

Performance issues

Bélády's anomaly: more memory, sometimes worse performance.

Principle of locality

- **Temporal locality**: Recently-accessed items are likely to be accessed in the near future.

- **Spatial locality:** Items whose addresses are near one another tend to be referenced close together in time.

So a page that is being accessed will have a **very high probability** to be accessed again.

More processes in the system → more CPU cycles are wasted. This is called **thrashing**.

