

# Do not distribute course material

You may not and may not allow others to reproduce or distribute lecture notes and course materials publicly whether or not a fee is charged.

# Topic 1

## Subcategory of Supervised Learning

### Linear Model

### Simple Linear Regression

A model assuming a linear relationship between the input variables and the output variable

Studied for over 200 years

Simple = only one input variable

CS/EE-UY 4563: INTRODUCTION TO MACHINE LEARNING  
PROF. LINDA SELLIE

Some of the slides are from Prof. Sundeep Rangan

Some approaches are taken from CMU 18-661 Introduction to Machine Learning

# Learning Objectives

---

- ❑ How to load data from a text file
- ❑ How to visualize data via a scatter plot
- ❑ Describe a linear model for data
  - Identify the label (target variable) and feature (predictor)
- ❑ Understand the objective function for linear regression
- ❑ Understand how partial derivatives can be used in a convex optimization problem
- ❑ Optimization:
  - Compute optimal parameters for the model using a closed form solution
  - Compute the optimal parameters for the model using gradient descent
- ❑ Evaluation:
  - Able to compute  $R^2$
  - Able to visually determine goodness of fit and identify different causes for poor fit

# Regression means real valued output

Used in statistics and Economics

---

THE TERM REGRESSION COMES FROM STATISTICS - WHENEVER HAVE A REAL VALUED OUTPUT WE CALL IT REGRESSION

# Notation

We will use the notation (mostly...) from Stanford and the Deep Learning Book

- Input (features):  $\mathbf{x} \in \mathbb{R}^d$  ( $\mathbf{x}^{(i)}$  for the  $i^{\text{th}}$  example)
- Output (target/label):  $y \in \mathbb{R}$  ( $y^{(i)}$  for the  $i^{\text{th}}$  example)
- Training data:  $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})$ ,

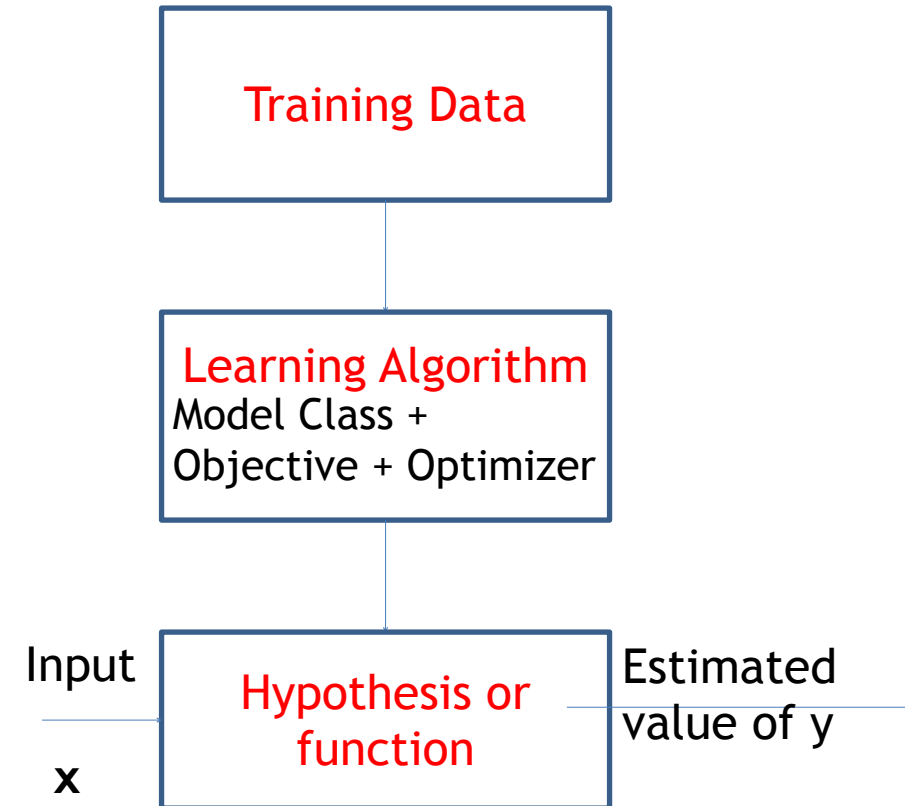
$$X = \begin{bmatrix} \mathbf{x}^{(1)T} \\ \mathbf{x}^{(2)T} \\ \vdots \\ \mathbf{x}^{(N)T} \end{bmatrix} \text{ design matrix}$$

- The number of training examples:  $N$

- Model class (hypothesis class):  $h : \mathbf{x} \rightarrow y$ , with  $h(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i$


$w_0, w_1$  are the *weights/parameters*

$w_0$  is the *bias*



# Outline

---

- 
- ❑ Motivating Example: Predicting the mpg of a car
  - ❑ Model: Linear Model
  - ❑ Objective function: Least Squares Fit Problem
  - ❑ Global Optimizer: LS Fit Solution
  - ❑ Local Optimizer: Gradient Descent
  - ❑ Assessing Goodness of Fit
  - ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation

# Predicting Trends

---

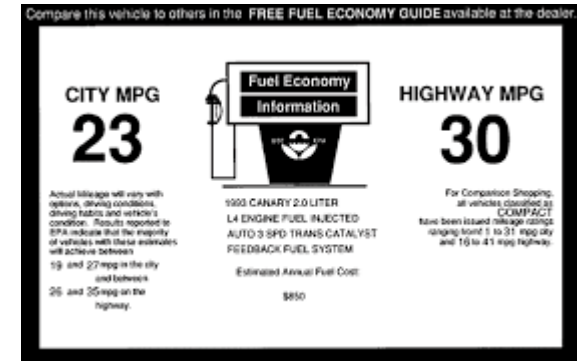
- ❑ Example: Predicting mpg for a car



mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
?	8	383.0	170.0	3563.0	10.0	70	1	dodge challenger se

# Example: What Determines mpg in a Car?

- ❑ What engine characteristics determine fuel efficiency?
- ❑ Why would a data scientist be hired to answer this question?
  - Not to help purchasing a specific car
    - The mpg for a currently available car is already known.
    - (If the car company isn't lying?)
  - To guide building new cars
    - Understand what is reasonably achievable before full design
  - To find cars that are outside the trend
    - Example: What cars give great mpg for the cost or size?





# Choose the average MPG?

---



# Keeping it simple

---

- Lets assume that one of the features can approximately predict the MPG



# Lets plot the data

---

# Demo on NYU Classes

## Simple Linear Regression for Automobile mpg Data

### Loading the Data

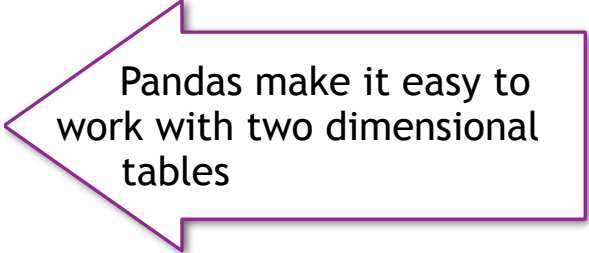
The python `pandas` library is a powerful package for data analysis. In this course, we will use a small portion of its features -- just reading and writing data from files. After reading the data, we will convert it to `numpy` for all numerical processing including running machine learning algorithms.

We begin by loading the packages.

```
In [86]: import pandas as pd  
import numpy as np
```

The data for this demo comes from a survey of cars to determine the relation of mpg to engine characteristics. The data can be found in the UCI library: <https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg>

You can directly read the data in the file, <https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data> We will load the data into ipython notebook, using the `pandas` library. Unfortunately, the file header does not include the names of the fields,



Pandas make it easy to work with two dimensional tables

The posted demo shows how to:






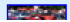
- Load data from a text file using the `pandas` package
- Create a scatter plot of data
- Handle missing data
- Fit a simple linear model
- Plot the linear fit with the test data
- Use a nonlinear transformation for an improved fit

# Getting Data

❏ Data from UCI dataset library: [UCI Machine Learning Repository](https://archive.ics.uci.edu/ml/datasets.html)

Browser: <https://archive.ics.uci.edu/ml/datasets.html?format=&task=reg&att=&area=&numAtt=&num>

Browse Through: 61 Data Sets Table View [List View](#)

Default Task - <a href="#">Undo</a>	Name	Data Types	Default Task	Attribute Types	# Instances	# Attributes	Year
Classification (257) <b>Regression (61)</b> Clustering (52) Other (51)	 <a href="#">3D Road Network (North Jutland, Denmark)</a>	Sequential, Text	Regression, Clustering	Real	434874	4	2013
Attribute Type Categorical (1) Numerical (52) Mixed (6)	 <a href="#">Air Quality</a>	Multivariate, Time-Series	Regression	Real	9358	15	2016
Data Type Multivariate (56) Univariate (4) Sequential (5) Time-Series (16) Text (4) Domain-Theory (3) Other (0)	 <a href="#">Airfoil Self-Noise</a>	Multivariate	Regression	Real	1503	6	2014
Area Life Sciences (9) Physical Sciences (8) CS / Engineering (25) Social Sciences (7)	 <a href="#">Amazon Access Samples</a>	Time-Series, Domain-Theory	Regression, Clustering, Causal-Discovery		30000	20000	2011
	 <a href="#">Auto MPG</a>	Multivariate	Regression	Categorical, Real	398	8	1993
	 <a href="#">Auto MPG</a>			Categorical			

← Today's lecture

# Loading the Data in Jupyter Notebook

## Try 1: The Wrong Way!

```
import pandas as pd
import numpy as np
```

```
In [67]: names = ['mpg', 'cylinders', 'displacement', 'horsepower',
                  'weight', 'acceleration', 'model year', 'origin', 'car name']
```

```
In [122]: df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data')
```

```
In [123]: df.head(6)
```

```
Out[123]:
```

	18.0 8 307.0 130.0 3504. 12.0 70 1 "chevrolet chevelle malibu"
0	15.0 8 350.0 165.0 3693. 11...
1	18.0 8 318.0 150.0 3436. 11...
2	16.0 8 304.0 150.0 3433. 12...
3	17.0 8 302.0 140.0 3449. 10...
4	15.0 8 429.0 198.0 4341. 10...
5	14.0 8 454.0 220.0 4354. 9...

- ❑ Python pandas library
  - `pd.read_csv` command
  - Reads URL or file location

- ❑ Creates a **dataframe** object
  - <http://pandas.pydata.org/pandas-docs/stable/dsintro.html#dataframe>

### ❑ Problems:

- Did not parse columns
  - All data in a single column
  - `read_csv` assumes columns are delimited by commas
- Mistakes first line as header

# Loading the Data in Jupyter

## Try 2: Fixing the Errors

```
In [125]: df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/'+
                        'auto-mpg/auto-mpg.data',
                        header=None,delim_whitespace=True,names=names,na_values='?')
```

You can display a first few lines of the dataframe by using head command:

```
In [126]: df.head(6)
```

Out[126]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18	8	307	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15	8	350	165	3693	11.5	70	1	buick skylark 320
2	18	8	318	150	3436	11.0	70	1	plymouth satellite
3	16	8	304	150	3433	12.0	70	1	amc rebel sst
4	17	8	302	140	3449	10.5	70	1	ford torino
5	15	8	429	198	4341	10.0	70	1	ford galaxie 500

- ❑ Fix the arguments in read\_csv
- ❑ pd.read\_csv has many options to deal with a wide variety of differently formatted data sets
- ❑ When you get a problem:
  - **Google is your friend!**
  - You are not the first to have these problems
  - Official documentation exists
- ❑ Ask questions on Brightspace if you get stuck

# Visualizing the Data using a scatter plot

```
In [150]: xstr = 'horsepower'
X = np.array(df[xstr])
y = np.array(df['mpg'])
```

```
In [146]: import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [151]: plt.plot(X,y,'o')
plt.xlabel(xstr)
plt.ylabel('mpg')
plt.grid(True)
```

- We will plot data in Python using Matplotlib
- A nice tutorial: [https://matplotlib.org/users/pyplot\\_tutorial.html](https://matplotlib.org/users/pyplot_tutorial.html)
- How could you predict the mpg of a car not in the data as a function of the horsepower? **Pair share**



$$(x^{(1)} = 130, y^{(1)} = 18)$$



$$(x^{(2)} = 165, y^{(2)} = 15)$$




$$(x^{(3)} = 150, y^{(3)} = 18)$$



$$(x^{(4)} = 150, y^{(4)} = 17)$$

...



$X =$

130
165
150
150
...

$y =$

18.0
15.0
18.0
16.0
...


$$x^{(1)} = 130 \quad x^{(2)} = 165$$

$$x^{(3)} = 150$$



# Outline

---

- ❑ Motivating Example: Predicting the mpg of a car
- ❑ Model: Linear Model
- ❑ Objective function: Least Squares Fit Problem
- ❑ Global Optimizer: LS Fit Solution
- ❑ Local Optimizer: Gradient Descent
- ❑ Assessing Goodness of Fit
- ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation

≈ approx.  
equal

# Approach

$$y \approx w_0 + w_1 \mathbf{x}$$


**Warning!** Stanford notes use  $\theta$   
instead of  $w$  for the parameters

$$\text{MPG} \approx w_0 + w_1 \text{Horsepower}$$

Learn parameters  $w_0, w_1$  of the green line

$$\text{Car 1: } w_0 + w_1 \cdot 130 = 18$$

$$\text{Car 2: } w_0 + w_1 \cdot 165 = 15$$



We are changing the notation from what you are used to seeing when we describe a line. Instead of  $mx + b = y$ , we are using  $w_0 + w_1 \mathbf{x}$

# Approach

$$y \approx w_0 + w_1 \mathbf{x}$$

$$\text{MPG} \approx w_0 + w_1 \text{Horsepower}$$


Learn parameters  $w_0, w_1$  of the green line

$$\text{Car 1: } w_0 + w_1 \cdot 130 = 18$$

$$\text{Car 2: } w_0 + w_1 \cdot 165 = 15$$

We can compactly represent this in matrix notation:

$$\begin{bmatrix} 1 & 130 \\ 1 & 165 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} 18 \\ 15 \end{bmatrix}$$



# Finding $w_0, w_1$ if we only use two examples

$$\begin{bmatrix} 1 & 130 \\ 1 & 165 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} 18 \\ 15 \end{bmatrix}$$

$$\begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \left( \begin{bmatrix} 1 & 130 \\ 1 & 165 \end{bmatrix} \right)^{-1} \begin{bmatrix} 18 \\ 15 \end{bmatrix}$$

$$= \begin{bmatrix} 4.71 & -3.71 \\ -.03 & .03 \end{bmatrix} \begin{bmatrix} 18 \\ 15 \end{bmatrix}$$

$$= \begin{bmatrix} 29.14 \\ -0.09 \end{bmatrix}$$

If we use more data, how can we estimate  $w_0, w_1$ ?

horsepower	mpg
130.0	18.0
165.0	15.0
150.0	18.0
150.0	16.0
...	...

Problem! There isn't a  $w_0, w_1$  that will satisfy all the equations

# Want to predict the best MPG based on the horsepower

$$\text{MPG} = w_0 + w_1 \text{horsepower} + \text{unexplainable\_stuff}$$

horsepower	mpg
130.0	18.0
165.0	15.0
150.0	18.0
150.0	16.0
...	...

Prediction car 1:  $w_0 + w_1 \cdot 130$

Prediction car 2:  $w_0 + w_1 \cdot 165$

Prediction car 3:  $w_0 + w_1 \cdot 150$

Prediction car 4:  $w_0 + w_1 \cdot 150$

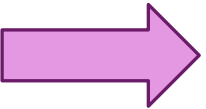
...

Learn  $w_0$  and  $w_1$

"Remember that all models are wrong; the practical question is how wrong do they have to be to not be useful." [George Box](#)

# Outline

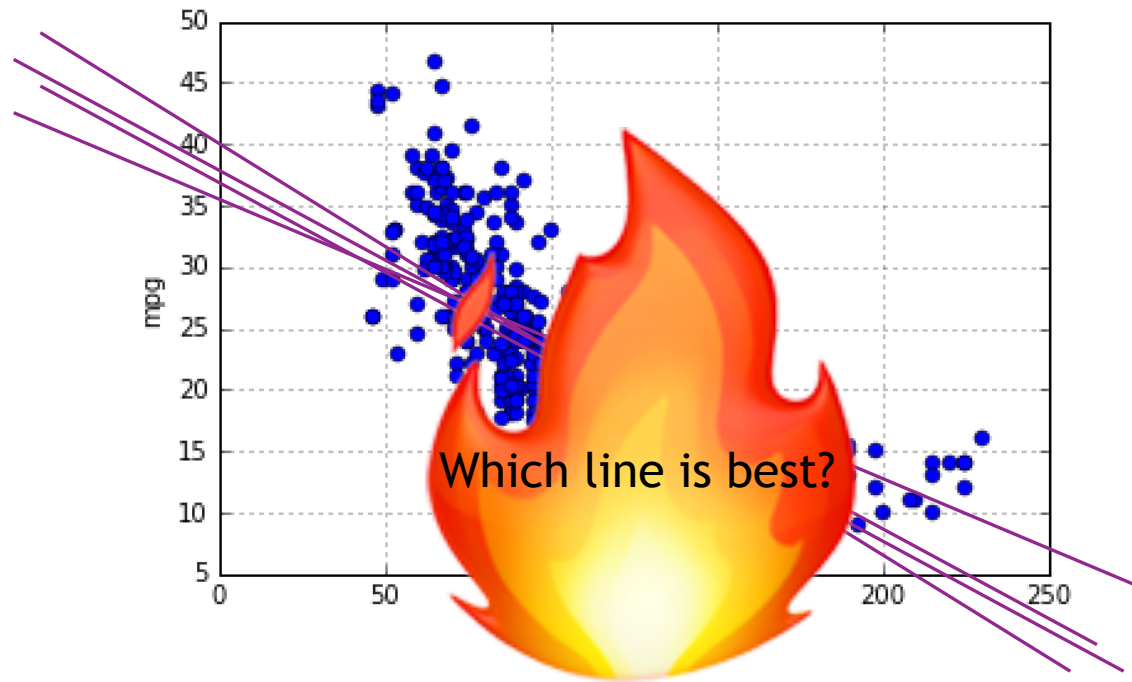
---

- ❑ Motivating Example: Predicting the mpg of a car
- ❑ Model: Linear Model
- ❑ Objective function: Least Squares Fit Problem
- ❑ Global Optimizer: LS Fit Solution
- ❑ Local Optimizer: Gradient Descent
- ❑ Assessing Goodness of Fit
- ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation



# Linear Model

---



# Pair share

## How do we decide which line is 'best'?

*What does it mean for one line to be better than another line?*

---


HOW MUCH DOES A POINT NOT ON THE LINE COST?

WE COULD CREATE A COST FUNCTION. THE BEST LINE HAS THE LOWEST COST WHEN SUMMED OVER ALL THE EXAMPLES

WHAT COST FUNCTION SHOULD WE USE?

# Outline

---

- ❑ Motivating Example: Predicting the mpg of a car
- ❑ Model: Linear Model
- ❑ Objective function: Least Squares Fit Problem
- ❑ Global Optimizer: LS Fit Solution
- ❑ Local Optimizer: Gradient Descent
- ❑ Assessing Goodness of Fit
- ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation

We choose to minimize the  
variation around the line

# How to measure errors?

## □ Absolute value of difference?

$$|y - \hat{y}| = |\text{MPG} - \text{predicted MPG}|$$

$$\begin{aligned} &|18 - (w_0 + w_1 * 130)| \\ &+ |15 - (w_0 + w_1 * 165)| \\ &+ |18 - (w_0 + w_1 * 150)| \\ &+ |16 - (w_0 + w_1 * 150)| \\ &\dots \end{aligned}$$

Note that

$$\begin{aligned} |y - \hat{y}| &= |\hat{y} - y| \\ \text{E.g. } |w_0 + w_1 * 130 - 18| \\ &= |18 - (w_0 + w_1 * 130)| \end{aligned}$$

## □ Squared value of difference?

$$(y - \hat{y})^2 = (\text{MPG} - \text{predicted MPG})^2$$

$$\begin{aligned} &(18 - (w_0 + w_1 * 130))^2 \\ &+ (15 - (w_0 + w_1 * 165))^2 \\ &+ (18 - (w_0 + w_1 * 150))^2 \\ &+ (16 - (w_0 + w_1 * 150))^2 \\ &\dots \end{aligned}$$

Note that  $(y - \hat{y})^2 = (\hat{y} - y)^2$

$$\begin{aligned} \text{E.g. } (w_0 + w_1 * 130 - 18)^2 \\ = (18 - (w_0 + w_1 * 130 - 18))^2 \end{aligned}$$

$$\text{mpg} \approx w_0 + w_1 \text{ horsepower}$$

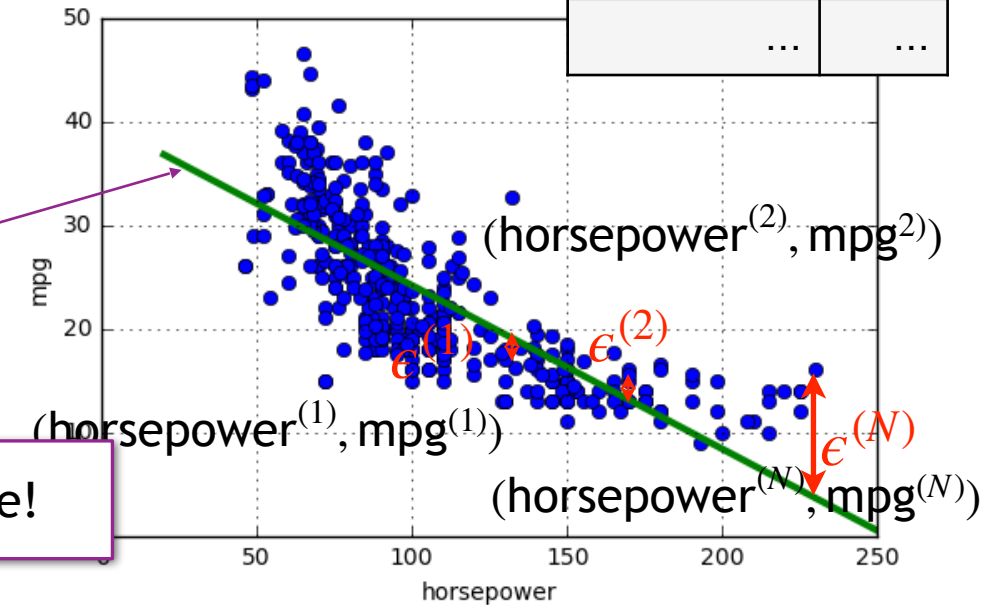
$$\hat{y} = w_0 + w_1 \mathbf{x}$$

$$\text{If } w_0 = 39.94, w_1 = -0.16$$

$$\hat{y} = h(\mathbf{x}) = 39.94 + (-0.16)\mathbf{x}$$

horsepower	mpg
130.0	18.0
165.0	15.0
150.0	18.0
150.0	16.0
...	...

Regression line



Differentiable!

# Least Squares Model Objective function

□ Model relationship between horsepower and mpg as a line  $\hat{y} = h(\mathbf{w}) = w_0 + w_1 \mathbf{x}$


□ Objective function: Find parameters  $\mathbf{w} = (w_0, w_1)^T$  to minimize cost

$$\text{RSS}(w_0, w_1) = \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 \mathbf{x}^{(i)}))^2 = \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2$$

Residual Sum of Squares  
Also called the **sum of squared residuals (SSR)** and **sum of squared errors (SSE)**

horsepower	mpg
130.0	18.0
165.0	15.0
150.0	18.0
150.0	16.0
...	...

$$\text{MSE}(w_0, w_1) = \frac{1}{N} \text{RSS}(w_0, w_1) = E_{\text{in}}(w_0, w_1)$$

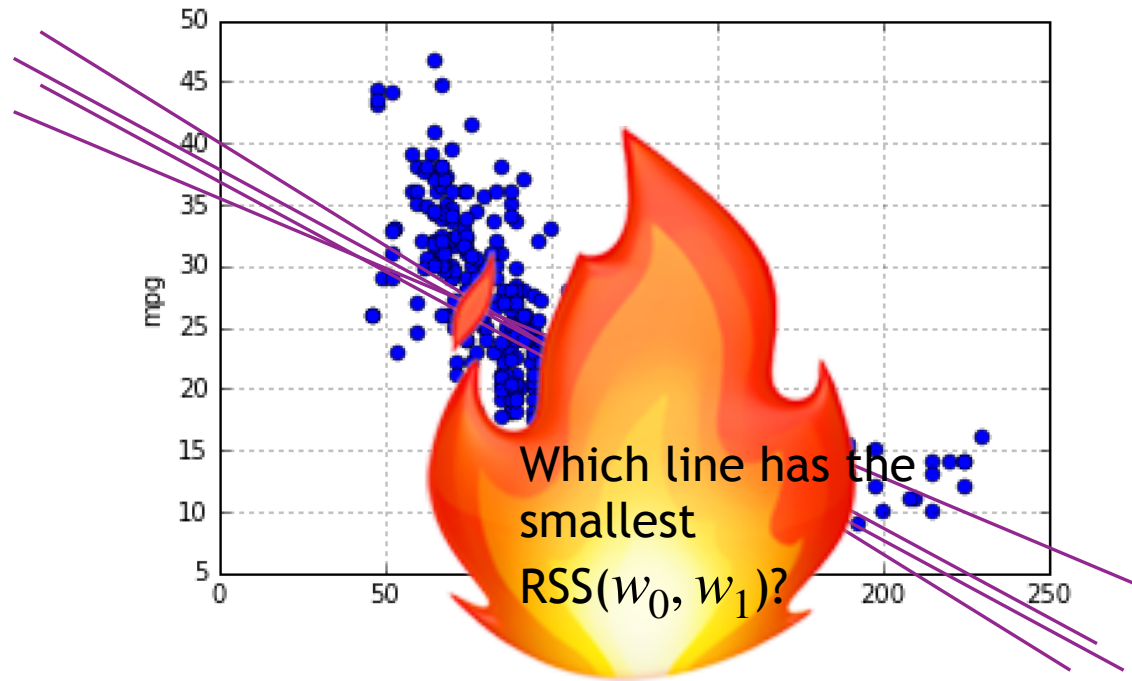



Regression line

$$\text{mpg} = w_0 + w_1 \text{horsepower}$$


# Linear Model

---






horsepower	mpg
130.0	18.0
165.0	15.0
150.0	18.0
150.0	16.0
...	...



$$\begin{aligned}
 &RSS(39.9, -0.2) \\
 &= (39.9 + (-0.2) * 130 - 18)^2 + (39.9 + (-0.2) * 165 - 15)^2 + \dots \\
 &= 18142.38
 \end{aligned}$$

$$\begin{aligned}
 &RSS(39.94, -0.16) \\
 &= (39.94 + (-0.16) * 130 - 18)^2 + (39.94 + (-0.16) * 165 - 15)^2 + \dots \\
 &= 9385.92
 \end{aligned}$$





# Finding Parameters via Optimization

## A general ML recipe

---

### General ML problem

- ✓ Get data
- ✓ Find a hypothesis class/model class with parameters
- ✓ Pick a loss function
  - Measures goodness of fit model to data
  - Function of the parameters

### Simple linear regression

→ Data:  $(\mathbf{x}^{(i)}, y^{(i)}), i = 1, \dots, N$

→ Linear model:  $\hat{y} = w_0 + w_1 \mathbf{x}$

→ Loss function:  $RSS(w) = \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2$

→ Find parameters that minimizes loss

→ Select  $w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$  to minimize loss function

Remember

$$\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$$


I.e. how do we  
find the line that  
*minimizes the cost  
function?*

# How do we find the global minimum?


$$\min_{\mathbf{w}} \text{RSS}(\mathbf{w})$$





Which values for  
 $w_0, w_1$  minimizes  
 $\text{RSS}(\mathbf{w})$ ?

# Convex function



# Not a convex function





# ...this would be even more difficult if we had more dimensions

---

I think we need an **algorithmic approach** to find the **optimal value**

**Mathematical optimization**

# Pair share

---

Discuss various methods for finding the optimal parameters, i.e. values for  $\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$


# Discussion of ideas

---

*Global optimization methods*

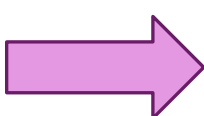
*Local optimization methods*

---



# Outline

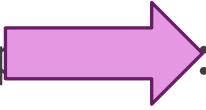
---

- ❑ Motivating Example: Predicting the mpg of a car
- ❑ Model: Linear Model
- ❑ Objective function: Least Squares Fit Problem
- ❑ Global Optimizer: LS Fit Solution
- ❑ Local Optimizer: Gradient Descent
- ❑ Assessing Goodness of Fit
- ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation



# Outline

---

- ❑ Motivating Example: Predicting the mpg of a car
- ❑ Model: Linear Model
- ❑ Objective function: Least Squares Fit Problem
- ❑ Global Optimizer : LS Fit Solution
- ❑ Local Optimizer: Gradient Descent
- ❑ Assessing Goodness of Fit
- ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation

# Global optimization for Linear regression

Given a dataset  $D = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$  how do we find the optimal weight vector  $\mathbf{w}$  to minimize

$$\text{RSS}(w_0, w_1) = \sum_{i=1}^N \left( y^{(i)} - (w_0 + w_1 \mathbf{x}^{(i)}) \right)^2$$

- We don't know the parameters  $w_0, w_1$  (they are also called coefficients)
- We could look over all possible lines to find the optimum...
- We can derive the partial derivative and then set the derivative to **zero** to find the **minimum value**
- Notice that since the error function is convex (we cannot get trapped at a local minimum - the solution is unique)

$$\frac{\partial \text{RSS}}{\partial w_0} = \sum_{i=1}^N 2(y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})(-1) = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) = \mathbf{0}$$

$$\frac{\partial \text{RSS}}{\partial w_1} = \sum_{i=1}^N 2(y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})(-\mathbf{x}^{(i)}) = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) \mathbf{x}^{(i)} = \mathbf{0}$$

$$\nabla \text{RSS}(w_0, w_1) = \begin{bmatrix} -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) \\ -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) \mathbf{x}^{(i)} \end{bmatrix}$$

# Method 1: Closed Form solution

$$RSS(w_0, w_1) = \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})^2$$

$$\frac{\partial RSS}{\partial w_0} = \sum_{i=1}^N 2(y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})(-1) = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) = 0$$

$$\frac{\partial RSS}{\partial w_1} = \sum_{i=1}^N 2(y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})(-\mathbf{x}^{(i)}) = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})\mathbf{x}^{(i)} = 0$$

Remember:

$$\frac{\partial (-a + b)^2}{\partial a} = 2(-a + b)^{2-1}(-1)$$

Simplify

$$\sum_i y^{(i)} = Nw_0 + w_1 \sum_i \mathbf{x}^{(i)}$$

$$\sum_i y^{(i)} \mathbf{x}^{(i)} = w_0 \sum_i \mathbf{x}^{(i)} + w_1 \sum_i \mathbf{x}^{(i)2}$$

Solving these we get the *least squares coefficient estimates* (posted slides will show the calculations)

$$w_0 = \bar{y} - w_1 \bar{x}$$

$$w_1 = \frac{\sum_{i=1}^N (\mathbf{x}^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^N (\mathbf{x}^{(i)} - \bar{x})^2}$$

$$\bar{y} = \frac{1}{N} \sum_i y^{(i)} \quad \bar{x} = \frac{1}{N} \sum_i \mathbf{x}^{(i)}$$

---

The next slide was not covered in class

✓ Take the partial derivatives

$$RSS(\mathbf{w}) = (\epsilon^{(1)})^2 + (\epsilon^{(2)})^2 + \dots + (\epsilon^{(N)})^2 = \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})^2$$

$$\frac{\partial RSS}{\partial w_0} = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)})$$

$$\frac{\partial RSS}{\partial w_1} = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) (\mathbf{x}^{(i)})$$

# Closed form solution

average sample mpg

average sample horsepower

✓ Set the partial derivatives to zero and solve for the parameter

I am going to remove this

$$\frac{\partial RSS}{\partial w_0} = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) = 0$$

$$Nw_0 = \sum_{i=1}^N y^{(i)} - \sum_{i=1}^N w_1 \mathbf{x}^{(i)}$$

$$w_0 = \frac{1}{N} \sum_{i=1}^N y^{(i)} - \frac{1}{N} w_1 \sum_{i=1}^N \mathbf{x}^{(i)} = \bar{y} - w_1 \bar{x}$$

$$\frac{\partial RSS}{\partial w_1} = -2 \sum_{i=1}^N (y^{(i)} - w_0 - w_1 \mathbf{x}^{(i)}) (\mathbf{x}^{(i)}) = 0$$

$$= \sum_{i=1}^N (y^{(i)} - \bar{y} + w_1 \bar{x} - w_1 \mathbf{x}^{(i)}) (\mathbf{x}^{(i)}) = \sum_{i=1}^N (\mathbf{x}^{(i)} y^{(i)} - \mathbf{x}^{(i)} \bar{y}) - w_1 \sum_{i=1}^N ((\mathbf{x}^{(i)})^2 - \mathbf{x}^{(i)} \bar{x})$$

$$w_1 = \frac{\sum_{i=1}^N (\mathbf{x}^{(i)} y^{(i)} - \mathbf{x}^{(i)} \bar{y})}{\sum_{i=1}^N ((\mathbf{x}^{(i)})^2 - \mathbf{x}^{(i)} \bar{x})}$$

$$\left( \sum_{i=1}^N \bar{x}^2 \right) - \left( \bar{x} \sum_{i=1}^N \mathbf{x}^{(i)} \right)$$

$$\left( \sum_{i=1}^N \bar{x} \bar{y} \right) - \left( \bar{x} \sum_{i=1}^N y^{(i)} \right)$$

□ Observe that:


$$\sum_{i=1}^N (\bar{x}^2 - \mathbf{x}^{(i)} \bar{x}) = 0$$

$$\sum_{i=1}^N (\bar{x} \bar{y} - y^{(i)} \bar{x}) = 0$$

$$w_1 = \frac{\sum_{i=1}^N (\mathbf{x}^{(i)} y^{(i)} - \mathbf{x}^{(i)} \bar{y}) + \sum_{i=1}^N (\bar{x} \bar{y} - y^{(i)} \bar{x})}{\sum_{i=1}^N ((\mathbf{x}^{(i)})^2 - \mathbf{x}^{(i)} \bar{x}) + \sum_{i=1}^N (\bar{x}^2 - \mathbf{x}^{(i)} \bar{x})}$$

$$= \frac{\sum_{i=1}^N (\mathbf{x}^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^N (\mathbf{x}^{(i)} - \bar{x})^2} = \frac{COV(X, y)}{VAR(X)} = \frac{S_{xy}}{S_{xx}}$$

this is the covariance



It has a closed form solution!

Note that  $\sum_{i=1}^N (\mathbf{x}^{(i)})^2 \neq \bar{x}^2$  while  $\sum_{i=1}^N \mathbf{x}^{(i)} = N \bar{x}$

# Code Notation

- ❑ The homework(lab = homework assignment)/Demos will use the following notation
- ❑ We will try to be consistent
- ❑ Note: Other texts use different notations

Notation			
Statistic	Notation	Formula	Python
Mean	$\bar{x}$	$\frac{1}{N} \sum_{i=1}^N x^{(i)}$	<code>xm = np.mean(x)</code>
(Population) variance	$s_x^2 = s_{xx}$	$\frac{1}{N} \sum_{i=1}^N (x^{(i)} - \bar{x})^2$	<code>Sxx = np.mean((x-xm)**2)</code>
(Population) covariance	$s_{xy}$	$\frac{1}{N} \sum_{i=1}^N (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})$	<code>Sxy = np.mean((x-xm)*(y-ym))</code>

# Auto Example

$$w_0 = \bar{y} - w_1 \bar{x}$$


$$w_1 = \frac{\sum_{i=1}^N (\mathbf{x}^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^N (\mathbf{x}^{(i)} - \bar{x})^2} = \frac{S_{xy}}{S_{xx}}$$

$X =$

130.0  
165.0  
150.0  
150.0  
140.0  
198.0  
220.0  
215.0  
...

$y =$

18.0  
15.0  
18.0  
16.0  
17.0  
15.0  
14.0  
14.0  
...



intercept =  $w_0 = 39.94$

slope =  $w_1 = -0.16$

```
xm = np.mean(X)
ym = np.mean(y)
```

```
syx = np.mean((y-ym)*(X-xm))
sxx = np.mean((X-xm)**2)
w1 = syx/sxx
w0 = ym - w1*xm
```

$$\bar{x} = (130 + 165 + 150 + \dots + 215)/392 = 104.47$$

$$\bar{y} = (18 + 15 + 18 + \dots + 14)/392 = 23.45$$

$$S_{xy} = ((130 - 104.47)(18 - 23.45) + \dots + (215 - 104.47)(14 - 23.45))/392 = -233.26$$

$$S_{xx} = ((130 - 104.47)^2 + (165 - 104.47)^2 + \dots + (215 - 104.47)^2)/392 = 1477.79$$

$$w_1 = (-233.26)/1477.79 = -0.16$$

$$w_0 = 23.45 - (-0.16)(104.47) = 39.94$$

```
yhat = w0 + w1*X
RSS = np.sum((yhat-y)**2)
MSE = np.mean((yhat-y)**2) # or RSS/N
```

$$RSS(39.94, -0.16) = 9385.92$$

$$MSE(39.94, -0.16) = 24.00$$

Calculations performed with higher precision

# Prediction


$$\hat{y} = h(\mathbf{x}) = w_0 + w_1 \mathbf{x}$$

$$w_0 = 39.94 \quad w_1 = -0.16$$

$$\text{MPG} \approx 39.94 + (-0.16) \text{Horsepower}$$

$$y \approx \hat{y} = 39.94 + (-0.16) \cdot 170 = 12.74$$


$$y \approx \hat{y} = 39.94 + (-0.16) \cdot 100 = 23.94$$






# Outline

---

- ❑ Motivating Example: Predicting the mpg of a car
- ❑ Model: Linear Model
- ❑ Objective function: Least Squares Fit Problem
- ❑ Global Optimizer: LS Fit Solution
- ❑ Local Optimizer: Gradient Descent
- ❑ Assessing Goodness of Fit
- ❑ Extra Slides: Global Optimizer for multivariate linear regression: Normal Equation

# Local Optimization





How can we make a small improvement  
in our parameters?

---

ANY IDEAS?

# Contour Plots

Graphical technique for representing a 3-D surface in 2-D plane



$$RSS(39.94, -0.16) = 9385.92$$

- **Contour curve:** a 3-D contour curve can be formed on the 3-D surface by finding all the points which have the same Z value (e.g.  $RSS(w_0, w_1) = k$ , where  $k$  is a constant value)
- **Level curve:** are in the 2-D plane and can be created by projecting the contour curve onto 2-D plane (e.g. the  $w_0, w_1$  axis)
- **Contour plot:** a 2-D plot containing level curves of a function