

8-9 Performance Analysis

Lecture 8-9 Performance Analysis

Standard Definition of Performance

For some program running on machine X,

$$\text{Performance}_X = 1 / \text{Execution time}_X$$

"X is n times faster than Y"

$$\text{Performance}_X / \text{Performance}_Y = n$$

Speedup

Number of cores = p

Serial run-time = T_{serial}

Parallel run-time = T_{parallel}

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$

Sources of Parallel Overheads

Overhead of creating threads/processes

Synchronization

Load imbalance

Communication

Extra computation

Memory access (for both sequential and parallel!)

Efficiency of a parallel program

$$E = \frac{\text{Speedup}}{p} = \frac{\frac{T_{\text{serial}}}{T_{\text{parallel}}}}{p} (\leq 1)$$

P = can be number of threads, processes, or core

Scalability

Scalability is the ability of a system to handle a growing amount of work efficiently.

If we keep the efficiency fixed by increasing the number of processes/threads and without increasing problem size, the problem is **strongly scalable**.

If we keep the efficiency fixed by increasing the problem size at the same rate as we increase the number of processes/threads, the problem is **weakly scalable**.

Execution Time

Elapsed Time (aka wall-clock time)

- counts everything (disk and memory accesses, I/O , etc.)
- a useful number, but often not good for comparison purposes

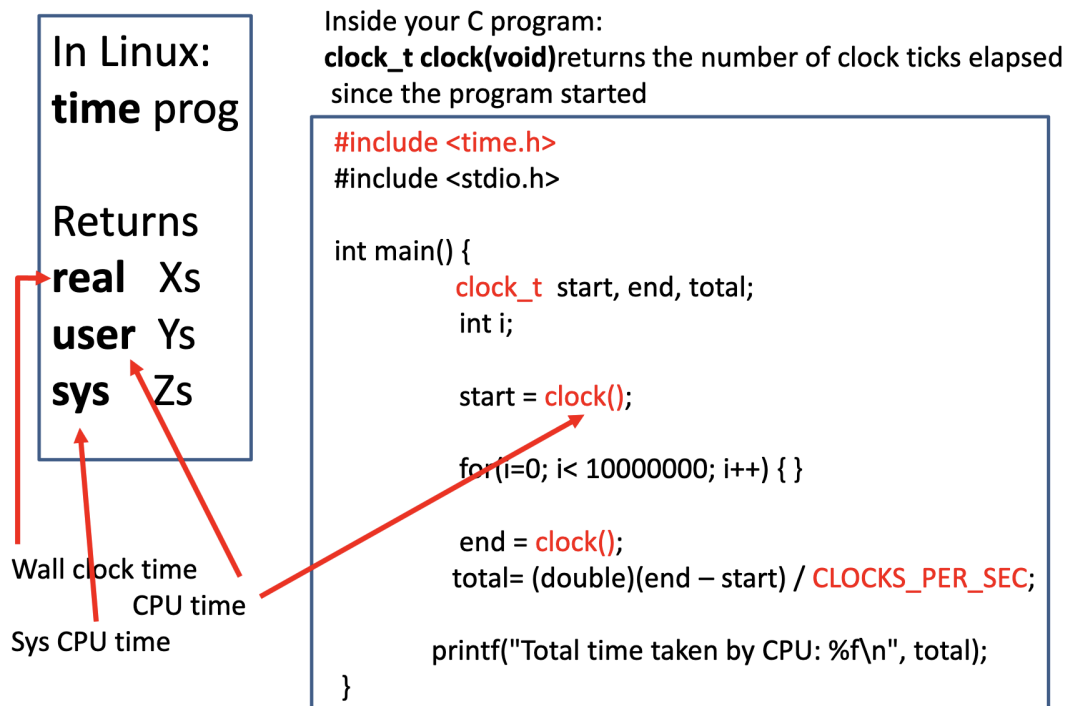
CPU time

- doesn't count I/O or time spent running other programs
- can be broken up into system time, and user time

Our focus: user CPU time

- time spent executing the lines of code that are "in" our program

Taking Timings



Simple Metrics

Response time (aka Execution Time)

- The time between the start and completion of a task
- e.g How long did the simulation take to finish?

Throughput

- Total amount of work done in a given time
- e.g. How many threads finished in the last two minutes?

Execution time for sequential program:

$$ET = IC * CPI * CT = \text{total_cycles} * CT = \text{total_cycles} / \text{clock_rate}$$

ET = Execution Time, IC = Instruction Count, CPI = Cycles Per Instruction, CT = Cycle Time

$$\text{MIPS} = \text{Million Instructions per Second} = \text{IC in millions} / \text{ET in seconds}$$

Pitfalls in timing in Parallel Machines

The total number of instructions executed may be different across different runs! (This effect increases with the number of cores.)

System-level code account for a significant fraction of the total execution time.

How to check the performance of a parallel machine?

Performance best determined by running a real application

Example of benchmarks: PARSEC (parallel), Rodinia (parallel), SPEC (sequential)

Conclusions

Execution time is what matters: system time, CPU time, I/O and memory time d

Scalability and efficiency measure the quality of your code.